





Learn R. By coding.

---

## **Copyright**

Title book: Learning R. By coding.

Author book: Thomas Kurnicki

© 2019, Thomas Kurnicki

Self publishing

ISBN 978-83-952046-1-6

ALL RIGHTS RESERVED. This book contains material protected under International and Federal Copyright Laws and Treaties. Any unauthorized reprint or use of this material is prohibited. No part of this book may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage and retrieval system without express written permission from the author / publisher

Dedicated to ...

Students...  
And anyone else who  
understands the importance of  
data.





# Table of Contents

Preface.....	1
1. R intro: the environment .....	3
1.1 Short passage about R .....	3
1.2 Installation.....	4
1.2.1 R Studio installation.....	5
1.2.2 Installation verification.....	6
1.2.2 Workflow in the R Studio environment. ....	6
1.3 Exercises .....	9
2. Data objects .....	12
2.1 Vectors, matrices, and data frames .....	12
2.2 Testing and changing the type of objects.....	16
2.3 Writing comments in your script .....	17
2.4 Sub-setting objects using names and indexes.....	19
2.5 Importing and exporting data sets.....	21
2.6 Cleaning and sub-setting your data .....	23
2.7 Exercises .....	24
3. Functions.....	28
3.1 Packages / libraries – installation and usage .....	28
3.2 Functions from libraries.....	30
3.3 User defined functions .....	35



3.4 Exercises .....	37
4. Loops.....	40
4.1 The “FOR” loop .....	42
4.2 The “WHILE” loop .....	45
4.3 Apply family function .....	47
4.4 Exercises .....	49
5. If statements .....	53
5.1 One possibility <i>if statement</i> .....	53
5.2 Nested <i>if statement</i> .....	54
5.3 Combining an <i>if statement</i> with a for loop .....	56
5.4 Exercises .....	57
6. Estimation and optmization .....	60
6.1 Linear regression .....	60
6.2 Testing for homoscedasticity. ....	64
6.3 Logistic regression .....	66
6.4 Predicting the dependent variable with a given model .....	71
6.5 Non-parametric estimation using smoothing splines ** .....	72
6.6 Fitting function parameters .....	75
6.7 Exercises .....	76

7. Data mining using SQL in R.....	80
7.1 The SQL template .....	80
7.2 The sqldf() function .....	83
7.3 SQL data joins .....	84
7.4 Exercises .....	88
8. Visualizations and interactive plots .....	93
8.1 The basic plot() function .....	93
8.2 GGLOT2 framework and library .....	95
8.3 Interactive visuals with Plotly .....	102
8.4 Exercises .....	107
9. Building a dashboard in R Shiny.....	111
9.1 Shiny as a framework .....	111
9.2 Server and UI interaction.....	112
9.3 Basic server functions .....	114
9.4 UI elements and user inputs .....	115
9.5 Exercises – R Shiny project.....	119
10. Report automation in R Markdown .....	121
10.1 R Markdown setup and basic features.....	121
10.2 Code chunks and types in Markdown .....	123
10.3 Exercises – R Shiny project ** .....	125
11. List of most useful functions.....	126
11.1 Mathematical .....	126

11.2 Data manipulation .....	126
11.3 Regression, optimization, and fitting .....	128
11.4 Descriptive statistics .....	128
12. Data mining with MongoDB (noSQL) .....	129
12.1 Installing MongoDB .....	129
12.2 Installing MongoDB Community with Homebrew	131
12.3 Running MongoDB.....	133
12.4 Using MongoDB from R .....	136
12.5 Uploading files to MongoDB.....	137
12.6 Managing non-relational data using NoSQL.....	137
12.7 Exercises .....	138
13. Text analytics .....	141
13.1 Importing data and creating a text corpus.....	141
13.2 Term-document matrix (TDM).....	143
13.3 Get sentiments from tidytext .....	145
13.4 Creating word clouds .....	147
13.5 Bayesian text classification model .....	149
13.6 Exercise.....	150
14. Appendix .....	151
14.1 READR cheat sheet [source: rstudio.com].....	152
14.2 GGLOT2 cheat sheet [source: rstudio.com].....	153





### **Preface**

First of all, I have to thank my wife for showing me a lot of support in writing this book. She always made references to the fact that learning a new coding language is just as hard as learning polish for a foreigner. This inspired me to help students practice R and create a repository of examples and exercises that boost the learning curve.

I'd also like to thank CC and XX. Both of them motivated me to study mathematics and enhance my coding skills. They've contributed many countless hours to my development.

This book is created for students and working professionals that are at the beginning of their journey with R. Prior experience or education in statistics or skills in another programming language is a "good-to-have" as this book will explain basic data manipulation, modeling, and programming frameworks. However, this book takes the reader from step one so that anyone, even without prior experience, can understand the content.

Each chapter will outline basic concepts and challenges and provide countless examples and exercises for each topic. It is highly recommended to follow the order of chapters as further chapters require understanding of the previous content. At the

## 1.1 Short passage about R

end of this book, the user should be able to create web applications and dashboards in Shiny.

Some readers might be disappointed due to the limited usage of buzz-words. “Machine learning, deep learning, and artificial intelligence” are avoided at all cost. Instead, the reader will receive a better understanding of basic programming concepts and will be exposed to more exercises and hands-on examples.

## **1. R intro: the environment**

### 1.1 Short passage about R

From the R project website (source: <https://www.r-project.org/about.html>):

*R provides a wide variety of statistical (linear and nonlinear modelling, classical statistical tests, time-series analysis, classification, clustering, ...) and graphical techniques, and is highly extensible. The S language is often the vehicle of choice for research in statistical methodology, and R provides an Open Source route to participation in that activity.*

*One of R's strengths is the ease with which well-designed publication-quality plots can be produced, including mathematical symbols and formulae where needed. Great care has been taken over the defaults for the minor design choices in graphics, but the user retains full control.*

*R is available as Free Software under the terms of the Free Software Foundation's GNU General Public License in source code form. It compiles and runs on a wide variety of UNIX platforms and similar systems (including FreeBSD and Linux), Windows and MacOS.*



## 1.2 Installation

As noted on the R-project website, R is the ultimate language for data scientists, statisticians, physicists, and many other professionals that need to work with data and want to make their lives easier.

### 1.2 Installation

In order to run R programs, you need to install the R compiler from the R-project website. The following URL works as of July 2018: <https://cran.r-project.org/mirrors.html>

This link has a list of all R mirrors. Find the location that is closest to your location and click on the URL. This will take you to the nearest CRAN (Comprehensive R Archive Network) and will list all available operating systems for which R is available. Download the one that will work with your operating system. After the download is complete, follow the prompts

If the link listed above doesn't work, use one of the popular search engines like [www.google.com](http://www.google.com) and search for R CRAN. The results should give you the most current URL under which the R CRAN is hosted.

### 1.2.1 R Studio installation

Even though having R from the CRAN is sufficient to run R programs on your local computer, you might want to consider using an IDE (Integrated Development Environment) to develop your code.

R Studio is so far, the best (according to the author) free IDE available today. It has many useful features such as:

- Matching special characters in your code, closing functions,
- Accessing all objects in the global and local environments,
- Browsing all plots created in the current R session.

R Studio is available from the following website:  
<https://www.rstudio.com>

There are a few versions of R Studio, including a commercial version for which the user has to pay. However, there is an open source version that is free and easy to download. All of the versions are available in the product section on the website.

### 1.2.2 Installation verification

For the first time you open R Studio, it will show 3 windows. We will explain the functionalities of all the windows in the next chapter.

In the big window on the left-hand side, called the ‘console’, type “Hello World” (you will notice that when opening the quotations, R Studio will automatically close them for you) and click enter. This will submit the “Hello World” statement.

This submission should result in the following response (printed in the next line):

```
[1] "Hello World"
```

If you see the above response, R Studio was installed successfully and you are ready to submit more complicated statements.

Also, you’ve just submitted your very first R code! Congratulations!

### 1.2.2 Workflow in the R Studio environment

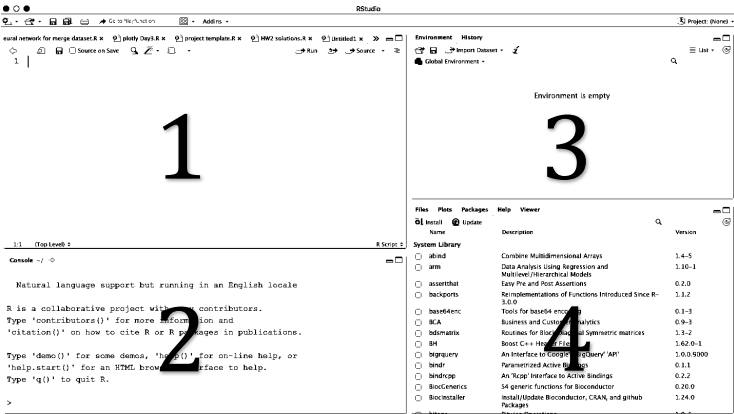
As mentioned in the previous chapter, for the first time you open R Studio you will see 3 windows:

- Console (left-hand side),
- Environment (upper right-hand side),
- Plots/packages/help window (lower right-hand side).

# 1.2.2 Workflow in the R Studio environment.

In addition to these 3 windows, it is very helpful to open a script that pops up as a 4<sup>th</sup> window in the upper left-hand side corner. In order to open a script window, go to the navigation panel at the top of your screen and click on the icon that looks like a white sheet of paper with a green “+” sign on it. When you get the drop down, select a new “script”.

At this point you should have 4 windows in your environment.



Picture 1.2.2.1: The R Studio environment [created by author]

Window number 1 is where you can browse all your scripts. It is possible to open more than one script if you need to reference different R programs. You will write your program in this window and submit it to the console using CTRL+ENTER. A good thing to point out at this point is that code is submitted line by line. Just

### 1.2.2 Workflow in the R Studio

environment.

place your cursor on the row you want to submit and use CTRL+ENTER. In order to run multiple lines at a time, select the desired lines using your mouse and when all are highlighted, use CTRL+ENTER.

Window number 2 is called the console. This is where all the statements you submit from the script are executed. I refer to this window as a dialog box with R, where the compiler “talks” with the user about any errors or results. Besides printing errors and objects, you can submit code in the console. However, you cannot save any of your work in the console to reuse it in your next R session.

In addition to the console, R Studio allows you to use a terminal in the same, lower left-hand side corner of your screen, just select the proper tab at the top of the 2nd window.

Window number 3 is called the environment. This is where you will see all objects created and saved in your R script. To create an object in your environment, you need to assign a name using the following template:

```
my_name <- my_object
```

The above template has three elements:

## 1.3 Exercises

- `my_name` – the name of the object, it will appear in the environment.
- `<-` the equal sign in R – this assigns the name to the object
- `my_object` – an object created using a function or by transforming an existing object

You can also use the “Import data” button to upload datasets in different formats. This feature is called the Data Import Wizard (DIW). When using the DIW, there will be a window in the bottom right hand corner of the screen where you can copy the code and reuse it in the future. What is more, as you progress in R coding you will discover the use of local and global environments.

Window number 4 has a lot of different applications. This window allows the user to:

- View and browse plots and graphs,
- Update, load, and install packages/libraries,
- Use the “help” feature to find information about libraries and functions.

### 1.3 Exercises

1. In the console, submit “Hi, my name is John”. What is the result?  
What was the impact on the other windows?

2. In a new script, submit using CTRL+ENTER:

```
"Hi, my name is Judy"
```

What is the result of this code? Were any of the other windows impacted?

3. Submit from the script using CTRL+ENTER:

```
?sd
```

What did you get in return? Can you tell what "sd" is and how to run it?

4. Submit from the script using CTRL+ENTER:

```
?sum
```

What did you get in return? Can you tell what "sum" is and how to run it?

5. Submit from the script using CTRL+ENTER:

```
?plot
```

Based on the help document, are the data inputs understandable? What options can we select for the plot?

6. Create an object in your environment called "my\_object" with the following content: "Hello master ". What data type was the new object assigned to?

7. Run the following:

```
testing_obj <- sd(c(1,2,3,4,5))
```

What was saved in the environment? How does this differ from the previous exercise?

8. Download the dataset from (copy and paste the link to your browser, the file will download automatically)  
[http://mba.tuck.dartmouth.edu/pages/faculty/ken.french/ftp/F-F\\_Research\\_Data\\_Factors\\_CSV.zip](http://mba.tuck.dartmouth.edu/pages/faculty/ken.french/ftp/F-F_Research_Data_Factors_CSV.zip)

and import the dataset using the DIW.

a) Copy the code from the DIW and paste it to your script.

b) Save your script on your local drive.

9. Open the file you downloaded to your local drive in #8 and delete the first column. Import the modified file using the script saved in 8.b



### **2. Data objects**

#### 2.1 Vectors, matrices, and data frames

Every object that you create in R is based on a vector. Even more complex objects are based on multiple vectors.

A vector is a sequence of numbers or characters in which order matters. It's best to imagine a simple vector of 2 elements. Let's consider a vector that describes a movement on a map. The first element (number) will describe how far an object moved in longitudinal distance. The second element (number) will tell us how far an object moved in the latitudinal distance. When we put these two numbers together, we get a vector describing how far the object moved in these two dimensions.

There are a few ways to create a vector (the following vectors create numeric vectors with integers):

```
X1 <- c(1,2,3)
X2 <- rep(c(1,2), each=4)
X3 <- seq( from=1, to=3, by=1)
```

The first approach (X1) is the most common, in which you can design any vector. Vector X1 has 3 elements: 1, 2, and 3. The

## 2.1 Vectors, matrices, and data frames

second method (X2) is used when we have a repetition of elements and the last method works best if there is linear pattern.

A vector can be classified by the type of data it contains.

Some of the basic vector types are (with examples):

- Numeric (listed in the code snippet above),
- Character - `c("January", "February", "March")`,
- Date - `c("2018-5-5", "2018-5-6")`,
- Factor - `c("yes", "yes", "yes", "no", "no")` but it translates to `c(1,1,1,0,0)`,
- Logical - `c("True", "False")`.

Stacking multiple vectors of the same type (numeric) will give us a matrix. A matrix will look like this:

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix}$$

This matrix would have [3,3] dimensions. The first number is the count of rows and the second number is the count of columns. In

## 2.1 Vectors, matrices, and data frames

this case we have 3 rows in 3 columns, thus the matrix dimensions are [3,3].

Here's a template:

```
my_object[row_index, column_index]
```

This template will be very useful in the next chapter, when we cover sub-setting matrixes and data frames. In order to create a new matrix use the `matrix()` function. The inputs are your numeric vectors.

A data frame looks very similar to a matrix. However, we need to mention that data frames have column and row names. What is more, all the data in one column needs to be of the same type.

```
> print(kyphosis)
  Kyphosis Age Number Start
1  absent  71      3     5
2  absent 158      3    14
3  present 128      4     5
4  absent   2      5     1
5  absent   1      4    15
6  absent   1      2    16
7  absent  61      2    17
8  absent  37      3    16
9  absent 113      2    16
10 present  59      6    12
```

This console screen shot presents the first 10 rows of the kyphosis dataset. Notice, that each column has only one data type. The Kyphosis column is a factor and all other columns are numeric. In this case, the row names are indexed from 1 to 10.

## 2.1 Vectors, matrices, and data frames

When referencing the columns, we use the data set name followed by the dollar sign. In this case it will be:

- `kyphosis$Age`,
- `kyphosis$Number`,
- `kyphosis$Start`,
- `kyphosis$Kyphosis`.

Please be careful with the Kyphosis column name, as the data set is called “kyphosis” (lower case K) and the column is called “Kyphosis” (upper case K). R is case sensitive.

The template for referencing a column is as follows:

```
dataset_name$column_name
```

Just like matrixes, data frames can be also referenced using indexes using template.

It is worth mentioning that there are other objects in R, such as user defined functions, lists, values, etc. Some libraries create library specific objects (xtm or xts for time series).

## 2.2 Testing and changing the type of objects

### 2.2 Testing and changing the type of objects

In the previous chapter, we've discovered different vector types. These types are assigned to vectors based on the type of data that they store.

The most powerful function to test an object (including a vector) for its type is:

```
typeof(object_name)
```

The user can also test an object for a certain type using a family of functions:

```
>is._type_(object_name)
```

Here are a few examples:

```
is.character(object_name)  
is.numeric(object_name)  
is.logical(object_name)  
is.data.frame(object_name)  
is.character(object_name)
```

For data frames and long vectors, there are a few functions that help the user understand the data. Most of these functions will give basic summary statistics and the type of data stored in the object:

```
summary(object_name)
```

## 2.3 Writing comments in your script

```
head(object_name)
str(object_name)
```

There are also many functions to change the type of data in an object. To convert data types use this family of functions:

```
as_type_(object_name)
```

Here are a few examples:

```
as.character(object_name)
as.numeric(object_name)
as.logical(object_name)
as.data.frame(object_name)
as.character(object_name)
```

### 2.3 Writing comments in your script

Before we cover sub-setting the objects, it is important to understand how to write comments and how they work in R.

Comments are extremely important for a few reasons. They create breaks in your code, so that you can make it more readable, you can explain what your code is doing, or you can comment out unused code for future reference without deleting it.

A comment starts with a “#”. When this special character is present, R skips further execution for a particular line.

## 2.3 Writing comments in your script

When “#” is the first character in a given line, the entire line is skipped by the R compiler.

```
# this will not be executed
```

However, if the “#” is after some code in the same line, the initial code will be executed:

```
my_object <- sd(1,2,3) # only the object will be executed
```

Remember to put as many comments in your code as possible. It'll make your life easier in the future.

It is extremely important to use comments when coding more advanced programs with loops or user defined functions. These structures can be lengthy and need to be opened with ( or { and closed with ) or }, depending on the structure. When there is more than a few lines of code inside these structures, it becomes almost impossible to find the closing ) or }. A good practice is to write comments that will tell the reader which loop or function is being opened or closed.

```
for( i in 1:100){          # opening the i-loop
  ...
  ...
  ...
}                          #closing the i-loop
```

## 2.4 Sub-setting objects using names and

indexes

### 2.4 Sub-setting objects using names and indexes

We've mentioned in the previous chapter that vectors are objects with elements for which order matters. This order is called indexing. Indexes start from 1 and are sequenced by 1.

```
my_vector <- c(2008, 2009, 2010, 2011)
indexes <- c(1,2,3,4) #or seq(from=1, to=4, by=1)
```

Sub-setting vectors can happen index by index or for a range of indexes.

Index by index (referencing data in my\_vector from the previous code snippet):

```
my_vector[1]
```

This will result in a one element vector with a value of 2008.

Range of indexes:

```
my_vector[1:3]
```

This will result in a vector with values of 2008, 2009, 2010.

Matrixes and data frames can be subset the same way. However, it is different from sub-setting vectors in the fact that they have two dimensions. Sub-setting needs to be performed on both, rows and columns:

```
new_df[row_index, column_index]
```



## 2.4 Sub-setting objects using names and

indexes

If one of the two indexes are empty, than all the elements pass on to the new object.

```
new_df[1:30, ] #all columns pass on to new object
```

As mentioned in chapter 2.1, data frames are objects with more attributes. The most important attribute of a data frame is its column names. We can subset one column using indexing (only one dimension as it becomes a vector):

```
mydata$mycolumn[1:30] # this will give us the first 30  
# elements of the column
```

If you want to create a new data frame with just a few columns, you can specify a vector with column names in the column index. Use the following template if you want to copy the “Date” and “Names” columns only:

```
my_new_df <- my_old_df[,c("Date", "Names")]
```

In order to subset a data frame by putting a filter on values in a given column, use the `which()` function:

```
sub_df <- mydf[which(mydf$Age>10), ]
```

In the example above, we have selected all the rows(observations) that have age above 10. The column index is

## 2.5 Importing and exporting data sets

empty, which means that all columns are transferred to the new data frame.

The `which()` function is also great when it comes to cleaning up any observation that has missing values in a given column(variable). It will need to be combined with the `is.na()` function:

```
sub_df <- mydf[-which(is.na(mydf$Age)), c("Age",  
"Number")]
```

The `which` statement pulls all the observations from the `Age` column (`mydf$Age`) that are `NA`s, in other words, have missing values. However, the main purpose was to eliminate them and this is why we used the minus sign in front of `which()`. The negative sign gave the remaining indexes, that were not part of the `which(is.na())` statement.

### 2.5 Importing and exporting data sets

You can import data using the Data Import Wizard that is available from the environment window. Another way to import a dataset would be to use one of the functions from the `readr` package. A `readr` cheat sheet is available in the Appendix.

## 2.5 Importing and exporting data sets

Before you write the function to import a dataset, it is highly recommended to convert the data to a text format .csv or an Excel format .xls (.xlsx works too).

The most commonly used template looks like this:

```
library(readr)
mydf <- read_csv("file.csv")
```

The *file.csv* argument has to be replaced with the file's directory. WARNING! You need to replace a forward slash "/" with a backslash "\" in the path as R requires backslashes when referencing a directory.

There is a similar function to write a .csv file containing an R data frame.

```
library(readr)
write_csv(x, "path.csv")
```

Where *x* is the data frame from the R environment and *path.csv* is the local directory where you want to write the file.

For more information on how to import or export data, go to the Appendix chapter and look for the "readr" cheat sheet.

### 2.6 Cleaning and sub-setting your data

After importing the data, open it using one of the following functions:

```
View(data_name)
print(data_name)
```

Do you understand all the variables (columns)? Do you need to go back to your metadata to learn more about your dataset? Are all the imported variables important?

Once you have a good understanding of the data, you can start sub-setting it. All the templates from the previous chapter are very helpful.

```
my_new_df <- my_old_df[,c("Date", "Names")]
sub_df <- mydf[which(mydf$Age>10), ]
sub_df1<- mydf[-which(is.na(mydf$Age)), c("Age",
"Number")]
```

There is also a wrapper function, easier to understand. We can rewrite the sub\_df data frame as follows:

```
sub_df <- subset(mydf, Age>10)
```

What is more, you might come across data that has a lot of missing values. In order to eliminate missing values from one column, you can use the **-which(is.na())** listed above. If there are multiple variables with missing values and you want to remove all observations that have at least one missing value, use the following:

```
clean <- na.omit(mydf)
```

### 2.7 Exercises

NOTE: Exercises marked with \*\* are advanced and require more time and research to solve.

1. Create vectors that have the following elements (do NOT use the `c()` function ):

`x1 = {10, 20, 30, 40, 50, 60, 70, 80, 90}`

`x2 = {60, 45, 30, 15, 0, -15}`

`x3 = {10, 100, 1000, 10000}`

`x4 = {243, 81, 9, 3}`

2. Create the following character vector and change it to numeric:

`y = {"1+2", "0101", "5*5"}`

## 2.7 Exercises

3. Import the kyphosis dataset using the library(rpart) and using the summary() function, give IQR (Inter Quartile Ranges) for the columns.

a) Change the \$Kyphosis column to a numeric variable (it will end up being binary)

b) Subset the dataset for the Age and Number columns. Create a new object called "my\_new\_kyphosis".

4. For exercises 1-3 write comments with an explanation of the logic you've used to create the objects.

5. Create a new vector with the following elements (use indexing):

- 1<sup>st</sup> element of vector x1 from exercise 1,
- 2<sup>nd</sup> element of vector x1 from exercise 1,
- 5<sup>th</sup> element of vector x2 from exercise 1,
- 1<sup>st</sup> element of vector x1 from exercise 1,
- 3<sup>rd</sup> element of vector x4 from exercise 2.

a) Subset the new vector for any elements that are greater or equal to 20.

6. Subset the my\_new\_kyphosis dataset, which was created in exercise 3, so that the result has all the observations with Age:

a) less than 80. Use the which() function

b) equal to 1. Use the which function. (hint: use "==" to compare)

c) export the data from letter “a” to your local drive.

7. Create the following matrix and get the diagonal using the sub-indexing frameworks provided in chapter 2.4:

10	11	9	15	19
52	19	7	10	22
28	40	6	99	33
35	26	5	87	91
0	12	16	81	200

8. Import the iris dataset (it's a build in R dataset).

- Create an object so that the dataset goes to your environment,
- Subset the entire dataset so that all the observations have Sepal Length greater than 5,
- What is the difference between the variable distributions (for all variables) of the entire iris dataset and the subset from 8b.

9\*\* Save a SAS dataset using the following data step:

```
>libname new "...your C://Desktop.... Path";
>data new.mysasdata ;
>input ID $ WEIGHT SIZE;
>dataline;
```

```
>101 150 3
>102 125 2
>105 210 6
>99 99 5
>110 199 5
> ;
>run;
```

- a) Import the SAS dataset using a function from the READR cheat sheet.
- b) What type are the columns?
- c) Are the column data types different from the data types in SAS?

10\*\*. Using the Fama French data from Exercise 8, Chapter 1, clean up all records that are missing. Make sure to exclude any missing values in all of the columns. Export your results to your local drive.

11\*\*. Find AQR factors from the [www.images.aqr.com/Insights/Datasets](http://www.images.aqr.com/Insights/Datasets) website and import the data directly from the url using the `download.file()` function. Is your data clean? If needed, subset and remove empty values from your dataset.



## 3.1 Packages / libraries – installation and usage

### **3. Functions**

When trying to solve any given problem or challenge using R, you will end up trying to find a function that solves your problem. Fortunately, any challenge that you might be facing was most likely encountered by someone else in the past and described or solved by other R community members. Since R is an open source project, people create functions and post them in packages on the CRAN. This is how you can solve your problem without reinventing the wheel. At this point you might be wondering how are you supposed to know which function to use for your problem. In this case a google search might be your best option.

### **3.1 Packages / libraries – installation and usage**

Packages (“libraries” is a synonym) are designed in a way to host multiple functions for a given subject. For example, if you needed to create multilayer plots (any type), you would call the ggplot2 package, or if you needed to build a decision tree, you would call the rpart package. Here’s a list of some of the most popular packages and the subject that they cover:

- ggplot2 – plotting,
- plyr – dataframe manipulation

### 3.1 Packages / libraries – installation

and usage

- stringr – character variable manipulation
- colorspace – defining colors
- reshape2 – data manipulation
- RColorBrewer – creating color palettes

All these packages can be downloaded from the CRAN using the

```
install.packages("name_of_package")
```

Other packages, that were developed recently, are located on GitHub. You will need to know the link in order to download these packages.

The installation process takes the file from CRAN or GitHub and installs it on your local drive, in a folder that R has created to store all packages. The `install.packages("")` function needs to be submitted only once. Next time, when you open R, all you will need to do is call the library without having to download it again.

A package/library is called using:

```
library(name_of_library)
```

You will need to call a library each time you open a new R session. To make a script more structured, it is recommended to list out

all the libraries, that the code requires, in the header of the script.

A script header may look like the example below:

```
#####  
#Created by XYZ on MM/DD/YYYY  
#Title of the script  
#####  
library(ggplot2)  
library(rprat)  
library(princomp)
```

### 3.2 Functions from libraries

As mentioned in the chapter 3.1, libraries combine multiple functions that cover the same subject. You can find information on each function in a library from its documentation. These documents are available online and can be found from the google search engine using key words such as “*library\_name* documentation pdf”. The description section of the document will tell us what the topic (subject) of all the functions is. The “R topic documented” section will give us a list of all functions that are in the rpart library. Each function will be described in the second part of the pdf document. The “Format” sections will list all the function inputs. It is very important to understand what objects and input formats the function needs to run properly.

Attached is a document that lists and describes all rpart library functions.

### Package ‘rpart’

February 23, 2018

**Priority** recommended

**Version** 4.1-13

**Date** 2018-02-23

**Description** Recursive partitioning for classification, regression and survival trees. An implementation of most of the functionality of the 1984 book by Breiman, Friedman, Olshen and Stone.

**Title** Recursive Partitioning and Regression Trees

**Depends** R (>= 2.15.0), graphics, stats, grDevices

**Suggests** survival

**License** GPL-2 | GPL-3

**LazyData** yes

**ByteCompile** yes

**NeedsCompilation** yes

**Author** Terry Therneau [aut],  
Beth Atkinson [aut, cre],  
Brian Ripley [trl] (producer of the initial R port, maintainer  
1999-2017)

**Maintainer** Beth Atkinson <atkinson@mayo.edu>

**Repository** CRAN

**Date/Publication** 2018-02-23 05:45:50 UTC

#### R topics documented:

car.test.frame . . . . .	2
car90 . . . . .	3
cu.summary . . . . .	5
kyphosis . . . . .	6
labels.rpart . . . . .	7
meanvar.rpart . . . . .	8
na.rpart . . . . .	9
path.rpart . . . . .	9

## 3.2 Functions from libraries

2

`car.test.frame`

<code>plot.rpart</code> . . . . .	10
<code>plotcp</code> . . . . .	12
<code>post.rpart</code> . . . . .	13
<code>predict.rpart</code> . . . . .	14
<code>print.rpart</code> . . . . .	16
<code>printcp</code> . . . . .	17
<code>prune.rpart</code> . . . . .	18
<code>residuals.rpart</code> . . . . .	19
<code>rpart</code> . . . . .	20
<code>rpart.control</code> . . . . .	22
<code>rpart.exp</code> . . . . .	23
<code>rpart.object</code> . . . . .	24
<code>rsq.rpart</code> . . . . .	25
<code>snip.rpart</code> . . . . .	26
<code>solder</code> . . . . .	27
<code>stagec</code> . . . . .	28
<code>summary.rpart</code> . . . . .	29
<code>text.rpart</code> . . . . .	30
<code>xpred.rpart</code> . . . . .	31

<b>Index</b>	<b>33</b>
--------------	-----------

---

<code>car.test.frame</code>	<i>Automobile Data from 'Consumer Reports' 1990</i>
-----------------------------	---

---

### Description

The `car.test.frame` data frame has 60 rows and 8 columns, giving data on makes of cars taken from the April, 1990 issue of *Consumer Reports*. This is part of a larger dataset, some columns of which are given in `cu.summary`.

### Usage

```
car.test.frame
```

### Format

This data frame contains the following columns:

**Price** a numeric vector giving the list price in US dollars of a standard model

**Country** of origin, a factor with levels 'France', 'Germany', 'Japan', 'Japan/USA', 'Korea', 'Mexico', 'Sweden' and 'USA'

**Reliability** a numeric vector coded 1 to 5.

**Mileage** fuel consumption miles per US gallon, as tested.

**Type** a factor with levels Compact Large Medium Small Sporty Van

**Weight** kerb weight in pounds.

**Disp.** the engine capacity (displacement) in litres.

**HP** the net horsepower of the vehicle.

car90

3

Source

*Consumer Reports*, April, 1990, pp. 235–288 quoted in  
John M. Chambers and Trevor J. Hastie eds. (1992) *Statistical Models in S*, Wadsworth and  
Brooks/Cole, Pacific Grove, CA, pp. 46–47.

See Also

car90, cu, summary

Examples

```
z.auto <- rpart(Mileage ~ Weight, car.test.frame)
summary(z.auto)
```

car90	Automobile Data from 'Consumer Reports' 1990
-------	--

Description

Data on 111 cars, taken from pages 235–255, 281–285 and 287–288 of the April 1990 *Consumer Reports Magazine*.

Usage

```
data(car90)
```

Format

The data frame contains the following columns

**Country** a factor giving the country in which the car was manufactured

**Disp** engine displacement in cubic inches

**Disp2** engine displacement in liters

**Eng.Rev** engine revolutions per mile, or engine speed at 60 mph

**Front.Hd** distance between the car's head-liner and the head of a 5 ft. 9 in. front seat passenger, in inches, as measured by CU

**Frt.Leg.Room** maximum front leg room, in inches, as measured by CU

**Frt.Shld** front shoulder room, in inches, as measured by CU

**Gear.Ratio** the overall gear ratio, high gear, for manual transmission

**Gear2** the overall gear ratio, high gear, for automatic transmission

**HP** net horsepower

**HP.revs** the red line—the maximum safe engine speed in rpm

**Height** height of car, in inches, as supplied by manufacturer

**Length** overall length, in inches, as supplied by manufacturer

Picture 3.2.1 Documentation for the rpart library. [source: R-CRAN]

Once we know which function we want to use, we can ask R to give us some help using the following template:

```
?function_name
```

For example, to get help on the `rpart()` function, we would submit the following (need to have the `rpart` package installed first):

```
?rpart
```

This will bring up the help document in the 4<sup>th</sup> window (lower right-hand corner) in our environment.

rpart {rpart}
R Documentation

## Recursive Partitioning and Regression Trees

**Description**

Fit a `rpart` model

**Usage**

```
rpart(formula, data, weights, subset, na.action = na.rpart, method,
      model = FALSE, x = FALSE, y = TRUE, parms, control, cost, ...)
```

**Arguments**

<b>formula</b>	a <u>formula</u> , with a response but no interaction terms. If this is a data frame, that is taken as the model frame (see <u><a href="#">model.frame</a></u> ).
<b>data</b>	an optional data frame in which to interpret the variables named in the formula.
<b>weights</b>	optional case weights.
<b>subset</b>	optional expression saying that only a subset of the rows of the data should be used in the fit.

Picture 3.2.2 Usage and arguments of the `rpart` function. [source: R-CRAN]

In the upper left corner, the help document will tell us the name of the library to which this function belongs. In this case it is `library(rpart)`.

The “Usage” section will have the function template. This is the most important section of the help document because it tells the

### 3.3 User defined functions

user how the inputs are named and what inputs the function needs.

The “Arguments” section will give some detail on each function input. Please keep an eye on the data types that are required and how the input object is supposed to be designed.

Many times, this document will provide a few examples on the last page for the user to reference.

#### 3.3 User defined functions

When trying to automate repetitive tasks, R allows to create user defined functions. A user created function must be saved to the environment with a custom name to be able to recall it later in the script. Here's the template:

```
function_name <- function( x , y, z) {  
    temp_object1 <- .....  
    temp_object2 <- .....  
    temp_object3 <- .....  
    return(temp_object3)  
}#this closes the function
```

The above script has a few elements:

- *function\_name* – this is the name of the function that will be used to recall it later in the script



### 3.3 User defined functions

- **function** – this is the function used to create a user defined function
- (x, y, z) – these are the function parameters a.k.a. inputs
- {...} - these French braces open and close the body of the function – this is where all the operations are saved
- **return(...)** – specifies which object is to be returned when the user defined function is called – usually a data frame, matrix, or vector – the return() statement is at the end of the body, before the closing braces

Calling the user defined function happens the same way as calling a function from a CRAN package. Using the above example, the call would look like this:

```
function_name(x=my_x, y=my_y, z=my_z)
```

This will return the temp\_object3 that is produced using the inputs that were specified in the function call.

One of the main reasons for using a user defined function is to avoid copy pasting the same code multiple times in one script when only a few objects or variables change each time. Instead, you can use a few function calls with different parameters.

```
output1 <- function_name(x=my_x1, y=my_y1, z=my_z)
output2 <- function_name(x=my_x2, y=my_y2, z=my_z)
...
```

```
output10 <- function_name(x=my_x10, y=my_y10,  
z=my_z)
```

### 3.4 Exercises

1. Install all most common packages on your local machine. You can find a list of packages in chapter 3.1
2. Create your own script header. Make sure you include all information about yourself so that it serves as a business card as well. Make sure it doesn't execute with the other R code using the "comment" option.
3. Find the best library to build a neural network (type: perceptron).
  - a) Install the library,
  - b) Read the library documentation from CRAN,
  - c) Get the R help document for the function that trains the model (hint: fist, find the function from the library document, secondly, use the "?" option to get help on the function)
4. Repeat exercise 3a,b,c for a library that allows you import financial instrument pricing data. For letter C, find the function that allows to import time series of prices for a given ticker.

### 3.4 Exercises

5. Get pricing data for the following tickers using the function from 4:

- a) SPY (the ETF that represents the market)
- b) S
- c) FB
- d) DB
- e) WFC

Save these in separate data frames in your environment.

6. Using the `Box.test` function (find the documentation to understand the function) test each of these time series datasets if they are stationary or non-stationary. (do you need to transform the data?)

- a) Which of these stocks can be modeled using ARMA and which using ARCH or GARCH?
- c) Build the ARMA or ARCH or GARCH models for each of these stocks.

7. Create a user defined function named *studentfunction* that transposes a numeric matrix (columns become rows) and subsets the matrix in a way so that only the first 5 rows remain in the dataset.

- a) Use the *studentfunction* function to transform the matrix created in Exercise 7 from chapter 2.

### 3.4 Exercises

b) Use the *studentfunction* function to transform the matrix that was given as an example in chapter 2.1

8. Create a user defined function named *transformmatrix* that takes the diagonal of a matrix and calculates a vector with two elements. Element one is the mean of the diagonal and element two is the median.

a) Use the *transformmatrix* function to transform the matrix created in Exercise 7 from chapter 2.

b) Use the *transformmatrix* function to transform the matrix that was given as an example in chapter 2.1

9\*\*. Create a user defined function that uses a *quandl* package to import zero coupon bond interest rates (the function inputs are the “from” and “to” dates and “N”) and calculates N-day moving averages for the following maturities:

a) 3M

b) 1Y

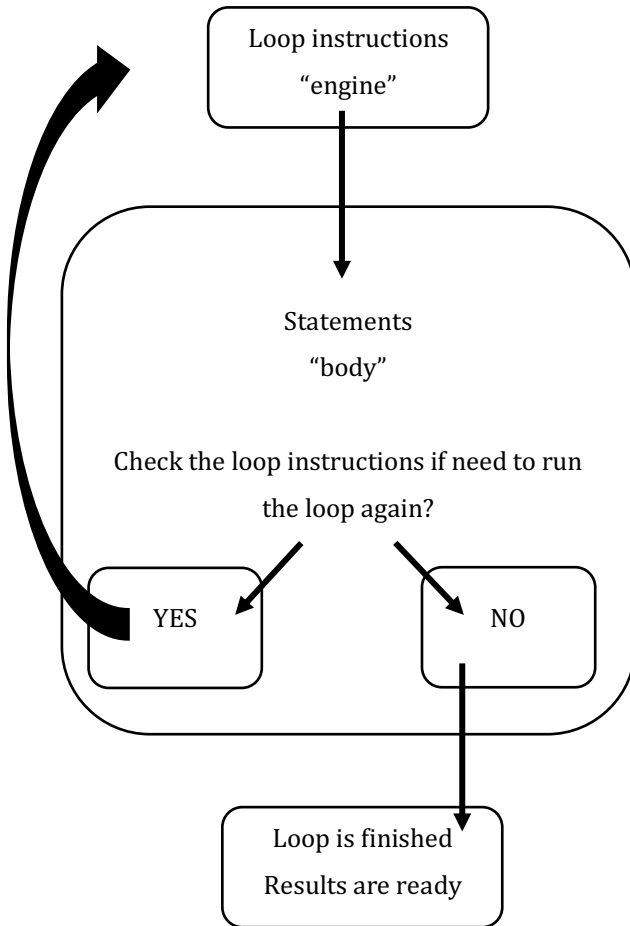
c) 10Y

d) 30Y

Consider  $N = \{25, 75, 250, 500\}$  ; from={"07-01-2011", "01-30-2010"} ; to={"12-20-2018", most\_currenet\_date -1 }

### **4. Loops**

Loops appear in almost any coding language. They are designed to run a certain chunk of code multiple times. Depending on the type of the loop, there might be conditions on how many times the loop runs.



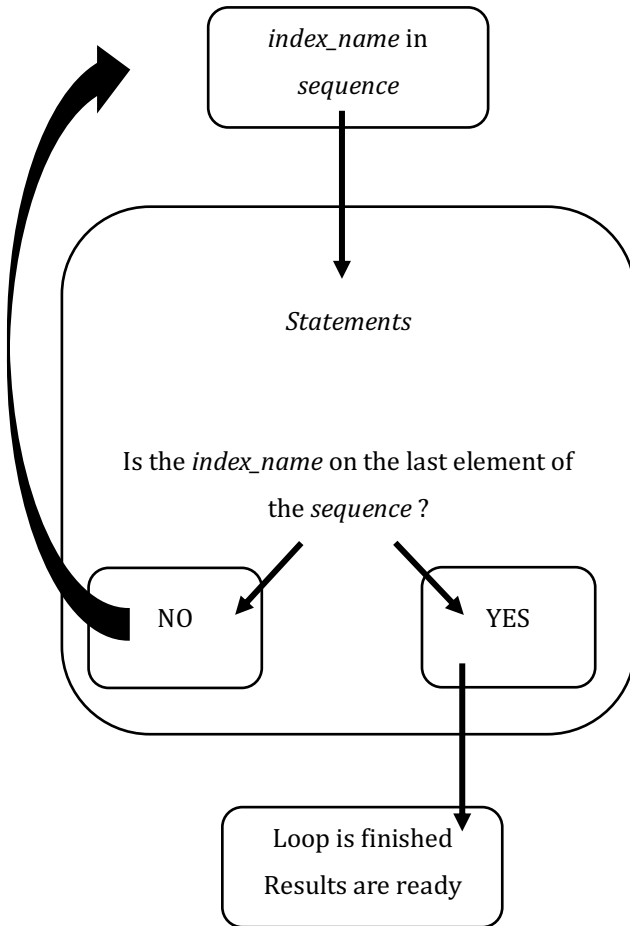
Graph 4.1: Generic loop design and data flow [designed by author].

### 4.1 The “FOR” loop

*For loops* run all the statements from the “body” for each element of a numeric vector. They are very powerful when it comes to manipulating and aggregating objects based on indexes. The template is as follows:

```
for( index_name in sequence) {  
    statements  
}
```

## 4.1 The “FOR” loop



Graph 4.1.1: *For loop* design and data flow [designed by author].

There are a few elements in this code chunk that need to be adjusted based on the data objects:



## 4.1 The “FOR” loop

- *Index\_name* – in most cases, the index name is a lower case letter such as i, m, or n.
- *Sequence* – this should resolve to a numeric vector with elements that should be used for creating the *index\_name*. A common solution is to write the following sequence with “:” such as 1:10. This would resolve to a vector (1, 2, 3, 4, 5, 6, 7, 8, 9, 10) and the loop would run 10 times.
- *statements* – (the “body”) is what is being submitted for each element in the *sequence*. Any *index\_name* reference will be substituted with a consecutive element from the *sequence*.

The first time the loop runs, the index name will take the first value from the sequence vector and substitute any reference in the statements. Each time the loop runs, the index name will take the next value from the sequence and put it in the statements. The loop stops when index name becomes the last element in sequence.

Let’s consider the following loop:

```
for( i in 1:2) {  
    print(i)  
}
```

This *for loop* will execute the following statements:

```
print(1)  
print(2)
```

## 4.2 The “WHILE” loop

The first time the loop runs, the “i” will be substituted with 1. The second time the loop runs, the “i” will be substituted with 2.

What is more, loops can be nested. For each run of the upper loop, the lower loop will run all the *sequence* elements.

```
for (n in 3:5) {                                #upper loop
    for( i in 1:2) {                            #lower loop
        print(i+n)
    }                                           #closing the i loop
}                                              #closing the n loop
```

This *for loop* will execute the following statements:

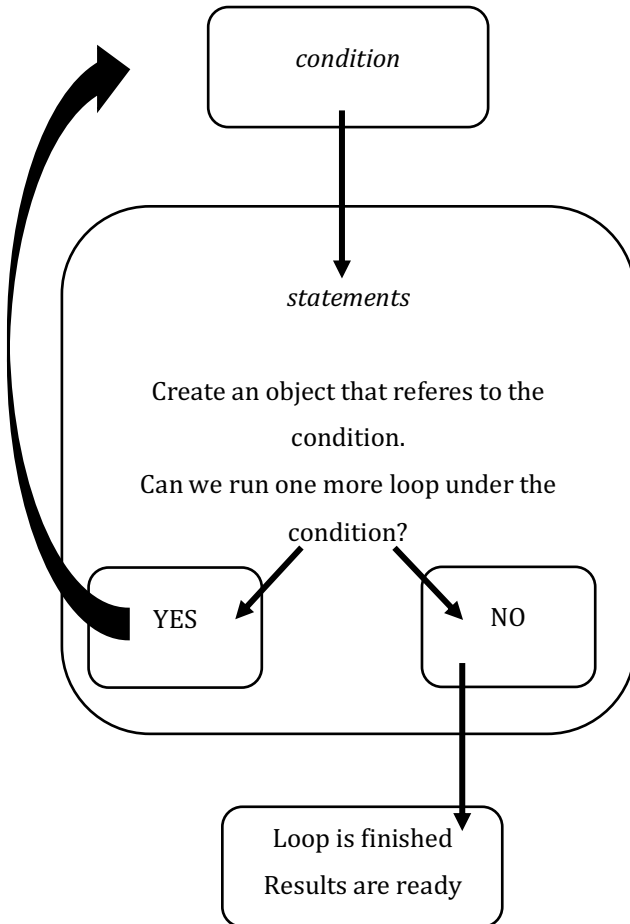
```
print(1+3)
print(2+3)
print(1+4)
print(2+4)
print(1+5)
print(2+5)
```

## 4.2 The “WHILE” loop

The while loop works in a similar way in which the *for loop* works. The only difference is that the while loop will stop executing when a condition is met.

## 4.2 The “WHILE” loop

```
while( condition ) {  
    statements  
    statement overwriting the object from condition  
}
```



Graph 4.2.1: While loop design and data flow [designed by author].

### 4.3 Apply family function

The condition should include a numerical object that has a starting value that meets the condition.

What is more, the numerical object from the *condition* should be overwritten as one of the last steps in the *statements* section.

Let's consider the following loop:

```
i <- 0 #assigning the initial value for i
while( i <3) {
  i <- i+1
  print(i)
}
```

This *for loop* will execute the following statements:

```
print(1)
print(2)
print(3)
```

### 4.3 Apply family function

Most *for loops* can be replaced by an “apply” function. The apply function takes a data object and performs a function on either the rows or columns (there are exceptions).

Here are all the functions from this family with their designated objects:

## 4.3 Apply family function

Function name	Data object
<b>apply()</b>	Columns or rows of a matrix
<b>lapply()</b>	Takes in a list and outputs a list
<b>sapply()</b>	Takes in a list and returns a vector
<b>tapply()</b>	Splits the input data based on a specified field
<b>mapply()</b>	Works the same way as sapply() but takes in multiple objects

Let's take a closer look at the `apply()` function as it is the most commonly used function from this family. When asking for help on this function (`?apply`), we get the following function description from CRAN:

### Usage

```
apply(X, MARGIN, FUN, ...)
```

### Arguments

**X** an array, including a matrix.

**MARGIN** a vector giving the subscripts which the function will be applied over. E.g., for a matrix 1 indicates rows, 2 indicates columns, `c(1, 2)` indicates rows and columns. Where `x` has named `dimnames`, it can be a character vector selecting dimension names.

**FUN** the function to be applied: see 'Details'. In the case of functions like `+`, `%*%`, etc., the function name must be backquoted or quoted.

**...** optional arguments to `FUN`.

Picture 4.3.1 Usage and arguments of the `apply` function. [source: R-CRAN]

The document presents the following template:

```
apply(x, margin, fun)
```

## 4.4 Exercises

The first argument, `x`, stands for the data object, a data frame or a matrix. The second argument, `margin`, ask for a numeric input, 1 or 2.

1- the function is applied to rows

2- the function is applied to columns

The 3<sup>rd</sup> argument, called `fun`, is expecting a function name that will be used to aggregate the data. This could be a mean, sum, sd, etc.

### 4.4 Exercises

1. For the matrix created in Exercise 7 from chapter 4:

- a) Calculate column medians using the `apply()` function.
- b) Calculate row means using the `apply()` function.

2. For the matrix that was given as an example in chapter 2.1:

- a) Calculate row medians using the `apply()` function.
- b) Calculate column means using the `apply()` function.

3. Recreate the same processes from Exercise 1 and 2 using a for loop.

4. Using stock tickers from Exercise 5 from chapter 3:

- a) Create a character vector with all tickers,

## 4.4 Exercises

b) Create a loop to perform the same test as in Exercise 6 (chapter#3).

5. Create a while loop that will create a numerical vector with a sequence from 13 to 63.

6. Create a nested for loop (upper loop accounts for row indexes and lower loop accounts for column indexes) that creates a data frame with TRUE or FALSE inputs (logical data frame). The TRUE or FALSE will be an outcome of testing each cell from the kyphosis dataset for numeric type. (Hint: use `is.numeric(kyphosis[upper_loop_index, lower_loop_index])`)

7. Create the same logical data frame as in Exercise 6 using a while loop.

8. Create a list of matrixes using the following function:

```
new_list <- replicate(10, matrix(rnorm(4), 2, 2))
```

a) Using the appropriate function from chapter 4.3, get the median for each row for each matrix in the `new_list`.

9. For the iris dataset (no need to call a function, iris is part of the base R) create for loop that does the following to each observation:

- a) changes the Species column from a character type to numeric. Assign 1 for setosa, 2 for virginica, and 3 for versicolor,
- b) creates a new column that groups the Petal.Length into 3 groups: group#1 for Petal.Length from 0 to 2, group #2 from 2.01 to 4.5, and group #3 from 4.51 to 7.

10. Create a user defined function that can clean up almost any data frame using a loop. The function inputs are the dataset name and the column indexes that we want to clean up. An example of the function call might look like this:

```
function_name(x=mydata, col_idx=c(1,2,3))
```

The loop inside the function will take column indexes (from the call inputs) and remove observations with empty values.

Hint: inside the loop, use:

```
new_df <- x[-which(is.na(x[,col_idx[i]]))]
```

Test your function with any dataset found on the [www.images.aqr.com/Insights/Datasets](http://www.images.aqr.com/Insights/Datasets) website.

11\*\*. Create a for loop that simulates independent Geometric Brownian Motion  $W_t$  ( $\sim N(0,1)$ ) for 5 hypothetical stocks.

$$\frac{dS_n}{S_n} = rdt + \sigma dW_t$$

The five stocks have the following annual growth rates and volatilities:

- a)  $r=0.01$  ,  $\sigma=0.1$
- b)  $r=0.02$ ,  $\sigma=0.12$



## 4.4 Exercises

c)  $r=0.03$ ,  $\sigma=0.14$

d)  $r=0.04$ ,  $\sigma=0.16$

e)  $r=0.05$ ,  $\sigma=0.18$

Simulate the portfolio return for the first two years, if each stock is equally weighted in the portfolio and the allocation doesn't change over time.

12\*\*. Create a for loop that creates a vector that creates elements using the following formula:  $x_{i+1} = \sum_{i=0}^i x_i$ .

Let  $x_0 = 3$  and  $x_1 = 5$ . Use  $i=10$  and  $i=20$ .

### **5. If statements**

The main purpose of using *if statements* is to create an action based on a logical statement. It is very similar to the “if” clause, type 1 conditional, used in English grammar:

IF -> this statement is true -> this will happen (execute in R).

#### **5.1 One possibility if statement**

When there is only one condition that needs to be checked in the *if statement*, the “one possibility” *if statement* can be used.

```
if (condition){outcome}
```

The *condition* should give a logical (TRUE/FALSE) answer. Here are a few statement examples:

- `if(mydf$Age[5] < 20){outcome}`
- `if(mydf$Name[1] == “Thomas”){outcome}`
- `if(my_constant == 10 ){outcome}`

Each time the *condition* is TRUE, then the outcome is executed. The outcome can be any action such as creating a new object, or overwriting a value in a data frame.

The examples above cover only the TRUE response. However, we can easily define an action when the statement is FALSE.

## 5.2 Nested if statement

```
if (condition){outcome when TRUE}else{outcome when FALSE}
```

In this case, the *if statement* will always produce an output.

Even though there is only one condition checked in these *if statements*, they can have a form of multiple sub-conditions. The most common way to link sub-conditions is with the “OR” or “AND” statements:

```
if (condition1 || condition2){outcome} # the || means OR  
if (condition1 && condition2){outcome} # the && means  
#AND
```

Example:

```
if(mydf$Age[5] < 20 && mydf$Name ==  
"Tom"){outcome}
```

### 5.2 Nested if statement

When there is a need for multiple condition statements, a nested *if statement* has to be used.

```
if (1st condition){outcome if 1st condition TRUE }  
else if (2nd condition){outcome if 2nd condition TURE}  
else {outcome if both are FALSE}
```

The above statement has a few elements:

- an opening *if statement* followed by the first condition
- multiple *else if statements* followed by different conditions
- an *else* statement that closes the entire string, it has the final outcome for a FALSE

All *condition* statements should give a logical (TRUE/FALSE) answer. Here are a few statement examples:

- `if(mydf$Age[5] < 20){outcome if TRUE}else{outcome if FALSE}`
- `if(mydf$Name[1] == "Thomas"){outcome if Thomas} else if (mydf$Name[1] == "Mary"){outcome if Mary} else {outcome if neither}`

The outcome can be any action such as creating a new object, or overwriting a value in a data frame.

Just like the one possibility *if statements*, the nested *if statements* can combine sub-conditions. The most common way to link sub-conditions is with the "OR" or "AND" statements:

```
if (condition1 || condition2){outcome} # the || means OR
if (condition1 && condition2){outcome} # the && means AND
```

## 5.3 Combining an if statement with a for

loop

### 5.3 Combining an *if statement* with a for loop

When using *if statements* on specific data points, we'll quickly realize that we need to test multiple data points in our data objects. For a vector, we might want to test all vector elements. For a data frame, the *if statement* can be applied for each cell in a given observation or variable. In these cases, we need to use the object's index(s) to write the *if statement*. The best way to run the "if" for each element is to put it in an indexing loop.

```
for ( i in 1:max_index){  
  
  if (1st condition with index i){outcome if 1st condition  
  TRUE } else if (2nd condition with index i){outcome if 2nd  
  condition TURE} else {outcome if both are FALSE}  
  
}#closing the for loop
```

The example below uses the "x" data frame, which has two columns. The first column is numeric and has a few "1" (observations with value 1). For each observation with a value 1 in col1, we want to put a value 0 in col2.

```

for(i in 1:nrow(x)){

  if(x$col1[i] == 1){x$col2[i] <- 0}else{x$col2[i] <- 1}

}#closing the for loop

```

### 5.4 Exercises

1. Considering the \$Age variable from the kyphosis dataset {library(rpart)} construct a one possibility *if statement* that:

- a) Creates a new column named \$group. For those observations that are older than 100 assigns a 1 and for those below or equal to 100, a 0 value.
- b) Creates a separate vector with a TRUE value for all that are older than 80, and a FALSE value for all that are younger or at 80.
- c) What is the type of the vector created in point b) . Can you convert it to numeric?

2. Considering the \$Number variable from the kyphosis dataset {library(rpart)} construct a nested *if statement* that:

- a) Replaces the current \$Number variable with the following names: “up to 3” for values 1,2,or3, “more than 3 up to 5” for values 4, and 5, “more than 5” for any values greater than 5.
- b) Creates a new column \$mynewcol with the following:

- for \$Kyphosis being “present” and \$Number above 4 create a “bad” label in \$mynewcol,
- for \$Kyphosis being “absent” and \$Number above 4 create a “good” label in \$mynewcol,
- for any other combination, create a “unsure” label in \$mynewcol.

3. Create a user defined function with an *if statement* that tests a vector from the function input. If the input vector is numeric, then the statement transforms the vector to character. If the input vector is character type, then the function returns the same character vector. The user defined function returns the new vector with all of its elements. Call the function using the following input vectors:

- a) `x1 <- c(1, 10, 5, 60, 80, 102, 101)`
- b) `x2 <- c(8, 1, 3, 5, 6, 100, 99, 98)`
- c) `x3 <- c("yes", "no", "yes", "no")`
- d) `x4 <- c("0101", "1", "2", "22", "022")`
- e) `x5 <- c("My", "favorite", "animal", "is", "a dog")`

4. Using the user defined function from Exercise 10, chapter 4, put the *for loop* into an *if statement*. The *if statement* will allow the loop to run only if the first column (specified in the `col_idx` input) has more than 25% missing values. Use the following functions to get the percentage of missing values.

```
>nrow(x[which(is.na(x[,i]))])/length(x[,i])
```

5\*\* Consider exercise 11 from chapter 4. As a last step, create an *if statement* that automatically excludes a stock from the final portfolio if it has lost value over the 2 years. If  $\frac{dS_{2years}}{S_0} < 0$  then exclude the stock and recalculate the weights for each remaining stock so that they are still equally weighted in the portfolio.



## **6. Estimation and optimization**

There are many types of regressions. However, all of them have one thing in common. They take independent variable(s) and estimate a dependent variable. The distribution of the dependent variable will dictate the type of estimation and type of regression that needs to be used. This chapter will cover the two most common estimation/regression methods outlined in the table below.

<b>Dependent variable distribution</b>	<b>Regression type</b>	<b>Estimation method</b>
<b>Continuous/Discrete</b>	Linear	min(RSS)
<b>Boolean/Binary</b>	Logistic	max(log likelihood)

Table 6.1 Linear and logistic regression characteristics [created by author].

Please keep in mind that there are other estimation methods that will work better with different data e.g. time series.

### **6.1 Linear regression**

Linear regression is a type of regression that fits a linear function to describe the relationship, called causation (not to be confused with correlation), between an independent variable(s) and the dependent variable. As stated in the prior chapter, the dependent variable should be continuous. Example of variables that can be regressed are: weight, height, length, size, salary, speed, age, etc.

## 6.1 Linear regression

Linear regression takes one or more independent variables  $X_1, X_2, X_3, \dots, X_n$  (usually represented by  $X$  with a lower index) as inputs. On the other side of the equation is the dependent variable  $Y$ . The regression will assign “weights” to each and every independent variable. These weights are called “estimated coefficients” and will be denoted by  $\widehat{\beta}_1, \widehat{\beta}_2, \widehat{\beta}_3, \dots, \widehat{\beta}_n$ .

For each set of beta, a new “estimated  $Y$ ” will be calculated, denoted by  $\hat{Y}$ . The linear regression formula is as follows:

$$\hat{Y} = \widehat{\beta}_0 + \widehat{\beta}_1 * X_1 + \widehat{\beta}_2 * X_2 + \widehat{\beta}_3 * X_3 + \dots + \widehat{\beta}_n * X_n$$

Where:

$\widehat{\beta}_0$ - is the intercept

$\widehat{\beta}_n$ - is the estimated n-th coefficient

$\hat{Y}$  - is the estimated dependent variable

There are different ways to estimate the best coefficients. The most popular method is called Ordinary Least Squares (OLS). The OLS minimizes the residual sum of squares (RSS) between the actual dependent response and the estimated value  $\hat{Y}$ .

In order to run a linear regression model in R, a data frame needs to be created with a few columns including the dependent and independent variables. The “dependent variable” column should

be a continuous numeric variable. A sample data frame can look like the one below (weight is our dependent variable).

weight	age	cars	salary	sex
199.5	50	2	150000	Male
180.0	20	0	90000	Male
198.9	53	1	130000	Female
150.3	39	1	100000	Female

In R, the most popular function to run linear regression is the `lm()` function.

```
lm(Y ~ X1+X2+X3+...+Xn , data=mydata)
```

lm (stats)

R Documentation

Fitting Linear Models

Description

lm is used to fit linear models. It can be used to carry out regression, single stratum analysis of variance and analysis of covariance (although `ov` may provide a more convenient interface for these).

Usage

```
lm(formula, data, subset, weights, na.action,
   method = "qr", model = TRUE, x = FALSE, y = FALSE, qr = TRUE,
   singular.ok = TRUE, contrasts = NULL, offset, ...)
```

Arguments

- formula** an object of class "formula" (or one that can be coerced to that class): a symbolic description of the model to be fitted. The details of model specification are given under 'Details'.
- data** an optional data frame, list or environment (or object coercible by `as.data.frame` to a data frame) containing the variables in the model. If not found in `data`, the variables are taken from `environment(formula)`, typically the environment from which `lm` is called.
- subset** an optional vector specifying a subset of observations to be used in the fitting process.
- weights** an optional vector of weights to be used in the fitting process. Should be `NULL` or a numeric vector. If non-`NULL`, weighted least squares is used with weights `weights` (that is, minimizing  $\sum(w_i * e_i^2)$ ); otherwise ordinary least squares is used. See also 'Details'.
- na.action** a function which indicates what should happen when the data contain `NA`s. The default is set by the `na.action` setting of `options`, and is `na.fail` if that is unset. The 'factory-fresh' default is `na.omit`. Another possible value is `NULL`, no action. Value `na.exclude` can be useful.

Picture 6.1.1 Usage and arguments of the `lm` function. [source: R-CRAN]

Y and X1:Xn are column names and do not need to be written with the data frame name and \$ sign because the data frame name is defined in the “data=mydata” argument.

The `lm()` function creates an object of class “lm” with a list of objects including all the model coefficients. Use the following code to get the coefficients:

```
my_coefficients <- mymodel$coefficients
```

It will also give the estimated “Y-hat” values:

```
my_estimates < mymodel$fitted.values
```

All the objects created by the `lm()` function are listed below.

### Value

`lm` returns an object of class “lm” or for multiple responses of class c(“mlm”, “lm”).

The functions `summary` and `anova` are used to obtain and print a summary and analysis of variance table of the results. The generic accessor functions `coefficients`, `effects`, `fitted.values` and `residuals` extract various useful features of the value returned by `lm`.

An object of class “lm” is a list containing at least the following components:

<code>coefficients</code>	a named vector of coefficients
<code>residuals</code>	the residuals, that is response minus fitted values.
<code>fitted.values</code>	the fitted mean values.
<code>rank</code>	the numeric rank of the fitted linear model.
<code>weights</code>	(only for weighted fits) the specified weights.
<code>df.residual</code>	the residual degrees of freedom.
<code>call</code>	the matched call.
<code>terms</code>	the <code>terms</code> object used.
<code>contrasts</code>	(only where relevant) the contrasts used.
<code>xlevels</code>	(only where relevant) a record of the levels of the factors used in fitting.
<code>offset</code>	the offset used (missing if none were used).
<code>y</code>	if requested, the response used.
<code>x</code>	if requested, the model matrix used.
<code>model</code>	if requested (the default), the model frame used.
<code>na.action</code>	(where relevant) information returned by <code>model.frame</code> on the special handling of NAs.

Picture 6.1.2 Values of the `lm` function. [source: R-CRAN]

## 6.2 Testing for homoscedasticity.

In chapter2, we've mentioned the `summary()` function. It is also very powerful when trying to summarize a model.

```
summary(model_name)
```

The `summary()` function will print all the coefficients, errors, and significance tests for each independent variable. This output might be helpful when assessing the quality of the model and whether any of the variables should be excluded. The p-value should give a good indication of the statistical significance of an estimate. (We want to get p-values that are smaller than 0.05 or 0.1, if bigger, the coefficient is statistically insignificant and needs to be equal to 0.)

The `summary()` function will also provide estimated values for coefficients  $\beta_n$ . The coefficient is easy to interpret because it describes the change in the dependent variable when the independent variable increases by one unit, *ceteris paribus*.

### 6.2 Testing for homoscedasticity.

Linear regression has a few assumptions, one of them being that the underlying data is homoscedastic. A good way of checking if the data is homoscedastic, is to create a scatter plot with the X variable on the X-axis and the Y variable on the Y-axis.

The chart below presents homoscedasticity and heteroscedasticity.

## 6.2 Testing for homoscedasticity.

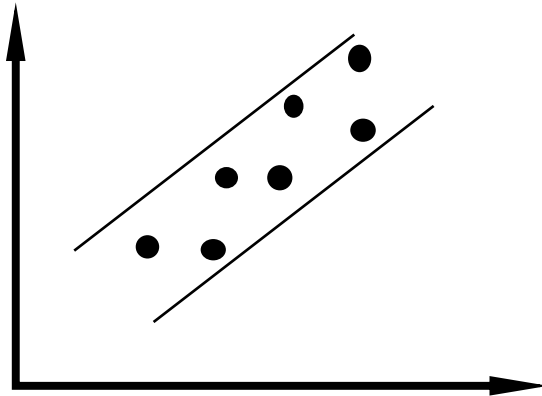


Chart 6.2.1 Homoscedasticity [created by author]

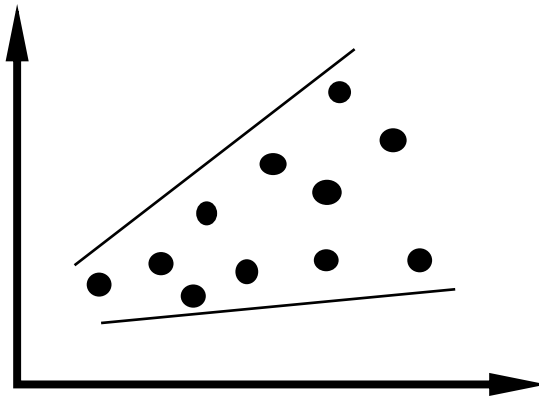


Chart 6.2.2 Heteroscedasticity [created by author]

## 6.3 Logistic regression

Homoscedasticity is observed when the variance of Y is constant over all the values of X, chart 6.2.1. When the variance of Y depends on the value of X (increases or decreases when X changes) then we say we have heteroscedasticity.

Besides creating a scatterplot using the `plot(x=my_x_variable, y=my_y_variable, type="p")` function, we can run the Breush – Pagan test.

```
library(lmtest)  
bptest(my_model_name)
```

If the p-value of this test is lower than the level of significance (lower than 0.05 or 0.1), you would reject the null hypothesis, and conclude that heteroscedasticity is present. The null hypothesis is that the data is homoscedastic (error variances are all equal). This is why we want bigger p-values to conclude homoscedasticity.

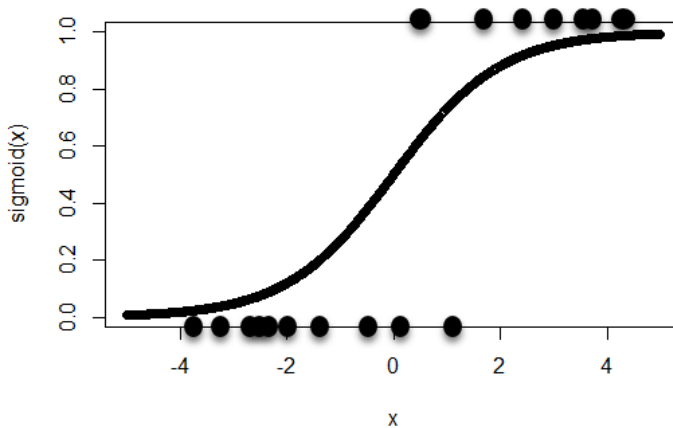
### 6.3 Logistic regression

Logistic regression is an estimation method that fits a sigmoid function between an independent variable and a binary variable. Binary variables are defined as variables with only two outcomes such as Yes/No, True/False, Male/Female, or 1/0. If the variable

## 6.3 Logistic regression

(vector in R language) is character or factor, than we should convert to a numeric type using any of the functions covered in previous chapters. After the conversion, the dependent variable should have only two values: 1 for the “success” outcome and 0 for the “failure” outcome.

The below chart shows a scatter plot (big black dots) representing our data. Notice, that the Y axis has only 2 values, 1 or 0.



Graph 6.3.1: Fitted sigmoid function for logistic regression (from author's collection)

The coefficient(s) of the sigmoid function is(are) fitted using the maximum likelihood estimation method. This function is extremely useful because it outputs the probability of success,



## 6.3 Logistic regression

denoted as  $p(1)$ , for any given value of  $X$ . With the sigmoid function, we can determine if an observation is more likely to be a 1 (if the probability is more than 0.5) or a 0 (if the probability is below 0.5).

As mentioned, the difference between a linear regression and logistic regression is that instead of estimating the value of the dependent variable, it will give the probability of 1,  $P(Y=1)$  for a binary variable (only 0 or 1 values). Here's the formula:

$$p(Y) = \frac{1}{1 + e^{-(\widehat{\beta}_0 + \widehat{\beta}_1 * X_1 + \dots + \widehat{\beta}_n * X_n)}}$$

Where:

$\widehat{\beta}_0$  - is the intercept

$\widehat{\beta}_n$  - is the estimated n-th coefficient

$p(Y)$  - is the probability of  $Y=1$

The data frame needs to be compiled the same way as for the `lm()` function. A sample data frame can look like the one below.

<b>purchase_decision</b>	<b>age</b>	<b>weight</b>	<b>salary</b>	<b>hair</b>
<b>1</b>	50	190	150000	Brown
<b>0</b>	20	150	90000	Brown
<b>1</b>	53	210	130000	Black
<b>0</b>	39	200	100000	Black

However, the function for a logistic regression is called `glm()` , with the `family="binomial"` option.

```
glm(Y ~ X1+X2+X3+...+Xn , data=mydata,
family="binomial")
```

Where `Y` and `X1:Xn` are column names and do not need to be written with the data frame name and `$` sign because the data frame name is defined in the `"data=mydata"` argument.

The `glm()` function will give the following output:

#### Value

`glm` returns an object of class inheriting from "glm" which inherits from the class "lm". See later in this section. If a non-standard method is used, the object will also inherit from the class (if any) returned by that function.

The function `summary` (i.e., `summary.glm`) can be used to obtain or print a summary of the results and the function `anova` (i.e., `anova.glm`) to produce an analysis of variance table.

The generic accessor functions `coefficients`, `effects`, `fitted.values` and `residuals` can be used to extract various useful features of the value returned by `glm`.

`weights` extracts a vector of weights, one for each case in the fit (after subsetting and na.action).

An object of class "glm" is a list containing at least the following components:

<code>coefficients</code>	a named vector of coefficients
<code>residuals</code>	the <i>working</i> residuals, that is the residuals in the final iteration of the IWLS fit. Since cases with zero weights are omitted, their working residuals are NA.
<code>fitted.values</code>	the fitted mean values, obtained by transforming the linear predictors by the inverse of the link function.
<code>rank</code>	the numeric rank of the fitted linear model.
<code>family</code>	the <i>family</i> object used.
<code>linear.predictors</code>	the linear fit on link scale.
<code>deviance</code>	up to a constant, minus twice the maximized log-likelihood. Where sensible, the constant is chosen so that a saturated model has deviance zero.
<code>aic</code>	A version of Akaike's <i>An Information Criterion</i> , minus twice the maximized log-likelihood plus twice the number of parameters, computed by the <code>aic</code> component of the family. For binomial and Poisson families the dispersion is fixed at one and the number of parameters is the number of coefficients. For gaussian, Gamma and inverse gaussian families the dispersion is estimated from the residual deviance, and the number of parameters is the number of coefficients plus one. For a gaussian family the MLE of the dispersion is used so this is a valid value of AIC, but for Gamma and inverse gaussian families it is not. For families fitted by quasi-likelihood the value is NA.

Picture 6.3.1 Values of the `glm` function [source: R-CRAN].

Just like with the `lm()` function, the `glm()` creates an object of class “glm” with a list of objects including all the model coefficients. Use the following code to get the coefficients:

```
my_coefficients <- mymodel$coefficients
```

The coefficients need to be transformed in order to drive business insight. We need to take the exponent of the coefficient value to make it understandable.

```
insight <- exp(my_coeff)
```

This exponent value can be interpreted as the change in odds ratio assuming a one unit change in the independent variable, *ceteris paribus*. E.g. the coefficient is 0.5 and  $\exp(0.5) = 1.64$  and after subtracting 1 we get 0.64. This value represents the change in odds ratio. In this case that would be an increase in odds ratio by 64%.

When we want to compare a few models (maybe with different independent variables) we can use the Akaike Information Criterion. The lower the AIC the better. The AIC can be obtained by the following:

```
myAIC <- mymodel$aic
```

## 6.4 Predicting the dependent variable with a given model

What is more, the `summary()` function works just the same way as it does for the linear model.

```
summary(mymodel)
```

The `summary()` contains all the most important information about the logistic regression model.

### 6.4 Predicting the dependent variable with a given model

The models that we've build in the chapters 6.1 and 6.2 can be used to predict the dependent variable for new data. The new data needs to have all the independent variables that were used to build the model (the column names in the new data frame have to be the same as the column names in the data frame used to build the model). There can be multiple observations in the new data frame.

To predict the dependent variable for a linear model, use the `predict()` function:

```
my_new_dependent <- predict(model_name, my_data)
```

Where:

*my\_new\_dependent* - will have all the predicted dependent outcomes,

## 6.5 Non-parametric estimation using smoothing splines \*\*

*model\_name* - is the name of the model that we've designed in chapter 6.1,

*my\_data* - is the name of the new data frame.

The function has one additional input when we want to predict the probability for new data using our logistic regression model:

```
new_prob <- predict(model_name, my_data,  
type="response")
```

The **type="response"** option will tell the predict() function that we need the true probability of 1,  $P(Y=1)$ . If the new data frame, *my\_data*, has more than one observation, the predict() function will output a vector of probabilities for each observation and save it in the "new\_prob" object.

### 6.5 Non-parametric estimation using smoothing splines \*\*

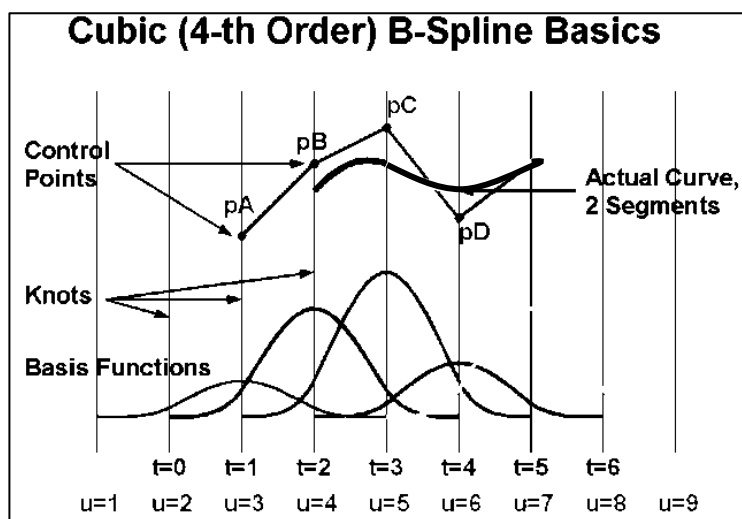
Not all data can be estimated using a linear regression especially when the relationship is non-linear. In this case, the best option would be to fit a n-th degree polynomial.

What is more, the relationship might change over time (usually X-axis as the predictor variable). This would mean that one polynomial cannot describe the entire time horizon.

## 6.5 Non-parametric estimation using smoothing splines \*\*

With such a problem, using a smooth spline would be the best option.

Splines work in such a way that they divide the time horizon (or other predictor variable) into smaller frames using knots. The curves between the knots can be estimated using different polynomials. This is why mathematicians use the term “piecewise polynomials” for splines.



Graph 6.4.1: Cubic spline framework [Retrieved on 8/22/2018 from: <https://people.eecs.berkeley.edu/~sequin/CS284>].

R has a function that will optimize all the knots and find the best degree to fit the polynomials on the data.

## 6.5 Non-parametric estimation using smoothing splines \*\*

```
my_spline_model <- smooth.spline(x=predictor,  
y=responses, df=number_df)
```

Where:

*x=predictor* – is the predictor column in a data frame e.g. time(as an index)

*y=responses* – this is the Y variable in your data frame

*df=number\_df* – the degree of polynomial that is being fitted, if not specified, the `smooth.spline()` will take the best fit.

Even though it's not required, it is a good practice to define the *df* in the `smooth.spline()` function. This way you can control the function and avoid overfitting. Many times, when the *df* option is omitted, the function will over fit the data.

The mathematical representation of the cubic spline function might be a bit confusing because the predictors are at higher degrees (n-degree polynomials). However, it is worth looking at it to have an understanding of what is being fitted on the data. Below is the function for cubic smooth spline regression:

$$E(Y|X) = \widehat{\beta}_0 + \widehat{\beta}_1 * X^1 + \widehat{\beta}_2 * X^2 + \widehat{\beta}_3 * X^3 + \widehat{\beta}_4 * (X - a_1)^3 + \\ + \widehat{\beta}_5 * (X - a_2)^3 + \widehat{\beta}_6 * (X - a_3)^3 + \dots + \widehat{\beta}_{k+3} * (X - a_k)^3$$

### 6.6 Fitting function parameters

Sometimes, there is a function or equation that is not standard and needs to be implemented via a user defined function (chapter 3.3). These have usually more parameters that need to be estimated. What is more, parameters are usually noted using Greek letters:  $\alpha, \beta, \gamma, \delta, \varepsilon, \theta$ , etc.

In this situation, there has to be a user defined function:

```
my_function <- function( $\alpha, \beta, \gamma, \delta, \varepsilon, \theta$ ) {  
    temp_object1 <- .....  
    temp_object2 <- .....  
    temp_object3 <- .....  
    return(temp_object3)  
}#this closes the function
```

And an optimization function that will optimize an objective considering the difference between the *temp\_object3* and real data. Usually, the objective is to minimize the RSS (residual sum of squares) but you can use a MLE (maximum likelihood estimation) to fit any nonlinear relationships. As you can probably imagine by now, the *temp\_object3*, which is returned by the function, and the *real\_data* object need to be vectors(or matrixes) of the same dimensions/length. The optimization



function comes from the library(minpack.lm) as has the following structure:

```
library(minpack.lm)
my_optimization <- nlsLM(real_data$Y ~
  my_function( $\alpha, \beta, \gamma$ ), start = c( $\alpha = 1, \beta = 1, \gamma = 2$ ),
  lower = c( $\alpha = -5, \beta = -10, \gamma = -12$ ),
  upper = c( $\alpha = 20, \beta = 15, \gamma = 7$ )
)
```

The `nlsLM()` will return a list of optimized values for the  $\alpha, \beta$ , and  $\gamma$  (function parameters). The “start” option defines the initial values at which the function will start the optimization, the “lower” will define a floor on each parameter, and the “upper” option will define a ceiling on the parameters.

Sometimes, there will be an error that says that the optimization hit a maximum number of allowed iteration. In this case, modifying the start, lower, or upper vectors might be helpful.

### 6.7 Exercises

1. Using the iris data frame (the data is part of base R and you can call it without installing any functions) build a linear regression that uses `Sepal.Length` as the dependent variable and the following as independent variables:

a) Sepal.Width

b) Petal.Length

c) Petal.Width

Build the 3 separate models with one independent variable per model. What can you say about the significance of the models and significance of the coefficients?

2. Using the kyphosis dataset (from library(rpart) ) build linear regression models with the Age variable being the dependent variable and the following as independent variables:

a) Number,

b) Start.

Build separate models for each independent variables and one multivariate regression model for the two variable together.

Are any of the models statistically significant? Are any of the independent variables insignificant ( $\alpha = 0.1$ )?

For the multivariate model, which variable contributes the most to the independent variable (based on the value of the estimate)?

3. Using the iris dataset:

a) combine the Setosa and Versicolor into group "0" and label the Virginica to "1". Create a new variable called iris\$Group with the 0 or 1 labels,

b) build a logistic regression model using any available data that will predict the observation being Virginica ( value of 1 in Group variable),

c) calculate the probability of a new plant being a Virginica for the following parameters:

Sepal.Width =5

Petal.Length =10

Petal.Width =7

Sepal.Length=9

4. Using the kyphosis dataset:

a) convert the kyphosis\$Kyphosis variable to numeric, assign a 1 to present and a 0 to absent,

b) build a logistic regression using all other variables and estimate the probability of the observation having a “present” hyphosis. What can you say about the coefficients? Are the significant?

c) calculate the probability of kyphosis being “present” for the following observation: Age=50, Start=10, Number=5.

5. Using all the single variable regressions from Exercise 1, test if the variable pairs are homoscedastic or heteroscedastic. Plot your findings. Using the `plot(x=my_x_variable, y=my_y_variable, type="p")` function. Use `my_data$variable_name` to define x and y variables in the function.

7. Using the Wage dataset from library(ILSR), use the `smooth.spline()` function to get a model for age being the independent variable and wage being the dependent variable? How good is your spline model?

8. Using the same function as in exercise 7, try to fit a cubic spline for age being the independent variable and logwage being the dependent variable (Wage dataset). Compare your results to the model from Exercise 7. Do they differ?

9\*\*. Fit the following polynomial function:

$$\hat{Y} = \hat{\beta} + \hat{\theta}X^{\delta}$$

on the cars dataset from R. Use speed as your independent variable X and dist (distance) as your dependent variable Y. Get the estimated coefficients using the `nlsLM()` function.

10. Test the following regressions for heteroscedasticity:

- from exercise 1 in this chapter, using the Breuch-Pagan test,
- from exercise 2 in this chapter, using the Breuch-Pagan test,
- from exercise 7 in this chapter, using a scatter plot [`plot()` function with `type="p"`].

## **7. Data mining using SQL in R**

### 7.1 The SQL template

The most basic SQL syntax includes 3 elements, the SELECT statement, the FROM statement, and the WHERE clause (the where clause is optional). The “SELECT statement” will list all the variable names that need to be exported from the table listed in the “FROM statement”.

```
SELECT variable_name1, variable_name2  
FROM my_dataset  
WHERE variable_name1 > value
```

If we wanted to get all the Number values and Kyphosis status for those observations that have Age greater than 100 from the kyphosis dataset, library(rpart) , we would use the following SQL query:

```
SELECT a.Number, a.Kyphosis  
FROM kyphosis a  
WHERE a.Age > 100
```

The lower-case letter “a” is a reference letter for the kyphosis dataset. Each variable name with “a.” in front of it, will be pulled

from the “a” table. For queries that use only one dataset, the reference letter can be omitted, but needs to be used when joining multiple tables (see chapter 7.3).

What is more, SQL offers basic aggregation functions. The aggregation can be a `SUM(variable_name)`, `COUNT(variable_name)`, `COUNT(DISTINCT variable_name)`, etc. For all of the aggregations, the SQL query needs to be extended with a `GROUP BY` clause. The `GROUP BY` will list the variables that are supposed to be combined at the aggregation level.

```
SELECT SUM(a.Number) AS Number_sum, a.Kyphosis
FROM kyphosis a
WHERE a.Age > 100
GROUP BY a.Kyphosis
```

A rule of thumb is to list all non-aggregated variables in the `GROUP BY` that are used in the `SELECT` statement.

Finally, SQL provides a very convenient and user intuitive way of creating labels and writing conditional statements. The `CASE WHEN..... THEN ..... ELSE .... END` is used very frequently and allows to perform similar operations as the “if statement” in R.

## 7.1 The SQL template

```
SELECT a.Kyphosis, CASE WHEN a.Number >20 THEN "older"  
ELSE "younger" END AS new_label  
FROM kyphosis a
```

The above SQL statement will create a new variable called `new_label` with "older" value when `Number >20` and "younger" label when `Number <=20`.

The AS statement after a variable aggregation or conditional statement assigns a new variable name. In the example above, the conditional statement was labeled as a new variable called "new\_label".

In the SELECT statement, each variable can be called separately, just like in the previous examples. However, all the variables can be called with the \*, just like in the example below:

```
SELECT *  
FROM kyphosis
```

This query will return all variables from the kyphosis dataset.

### 7.2 The sqldf() function

To run SQL in R, you need to install the SQLDF package and call the `library(sqldf)` in every new R session.

```
install.packages("sqldf")  
library(sqldf)
```

After calling the `library(sqldf)`, the `sqldf()` function will become available. It has the following structure:

```
new_df <- sqldf("  
                SELECT *  
                FROM kyphosis  
                ") #closing the sqldf function
```

This function requires the SQL statement to be surrounded by "" so that it appears as a character. What is more, the dataset used in the `sqldf()` function must be a data frame. The function will not work on a matrix, vector, or list. It is recommended to use the `as.data.frame()` function on the object to make sure it is a pure data frame before using the `sqldf()`.

In order to run the CASE WHEN example from the previous chapter, the following code has to be submitted:



```
new_data <- sqldf("
    SELECT a.Kyphosis,
    CASE WHEN a.Number >20 THEN "older"
    ELSE "younger"
    END AS new_label
    FROM kyphosis a
")
```

### 7.3 SQL data joins

When two or more data tables are in a relationship, they can be joined together. Both datasets will have one or more columns in common (that are the same). These columns are called primary key (for the table with distinct values in the common variable) and secondary key (for the table with multiple values in the common variable).

The relationship can be:

- one to one - one record from the first table corresponds to one record in the second table,
- one to many - one record from the first table corresponds to multiple records from the second table,
- many to one - multiple records from the first table correspond to one record from the second table,

### 7.3 SQL data joins

- many to many – many records from the first table correspond to multiple records from the second table.

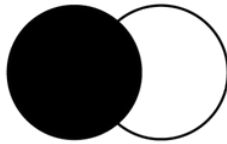
Once we know what the relationship looks like, we can determine which data join to use. What is more, the use of a particular data join will be determined by which data table is more important and if we want to exclude records from the less important data table.

There are multiple data joins that allow combining data in different ways. Here are the most common data joins:

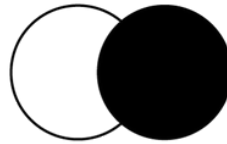
- INNER JOIN – keeps only the records that find a match from both tables,
- FULL OUTER JOIN – keeps all the records from both datasets,
- LEFT JOIN – keeps all the records from the “left” table and drops records from the “right” table that weren’t matched with the left table,
- RIGHT JOIN – keeps all the records from the “right” table and drops records from the “left” table that weren’t matched with the right table.

## 7.3 SQL data joins

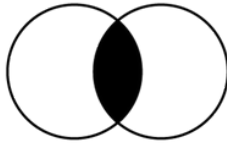
The chart below shows what data is being passed to the outcome table in each join. The black area represents the data that is passed and the white area is the data that is dropped:



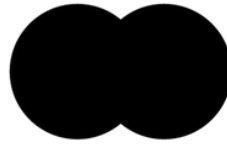
Left Join



Right Join



Inner Join



Full outer Join

Graph 7.3.1 Data join types. [Retrieved on 8/22/2018 from:

<https://www.quora.com/What-is-the-difference-between-joining-and-blending-in-Tableau>]

Each join has its own template. The most common join templates are provided below.

```
SELECT *  
FROM table1 a  
INNER JOIN table2 b  
ON a.key_variable=b.key_variable
```

```
SELECT *  
FROM table1 a  
LEFT JOIN table2 b  
ON a.key_variable=b.key_variable
```

```
SELECT *  
FROM table1 a  
RIGTH JOIN table2 b  
ON a.key_variable=b.key_variable
```

```
SELECT *  
FROM table1 a  
FULL OUTER JOIN table2 b  
ON a.key_variable=b.key_variable
```

Each of the templates above has a JOIN statement and an ON statement. The JOIN statement will indicate the type of join and the table name that is being joined. The ON statement indicates which variables are supposed to be used to create the join. These ON statement variables will be the primary key and secondary key.

The library(dplyr) has an easy to use family of functions to perform joins:

```
library(dplyr)
my_join <- left_join(dataset1, dataset2, by =
c('primary_key'))
my_join1 <- inner_join(dataset1, dataset2)
```

## 7.4 Exercises

1. Write a SQL query on the kyphosis dataset using the `squidf()` function to get the following:

- a) All the variables and all the observations,
- b) Just the Number and Age variables and all the observations,
- c) All the variables and only the observations that have Age greater than 50,
- d) Just the Number and Age variables and only the observations that have Age smaller than 40,
- e) Just the Kyphosis variable for all the observations with Number smaller than 5.

2. Write a SQL query on the iris dataset using the `squidf()` function to get the following:

- a) All the variables and all the observations,

- b) Just the Sepal.Length and Sepal.Width variables and all the observations ,
- c) All the variables and only the observations that have Sepal.Length greater than 4,
- d) Just the Pedal.Length and Sepal.Width variables and only the observations that have Sepal.Length smaller than 4,
- e) Just the Sepal.Length variable for all the observations with Sepal.Width smaller than 4.

3. Write a SQL query on the kyphosis dataset using the sqldf() function to get the specified aggregation for:

- a) All the observations. Aggregate the data to get the sum of Age per each kyphosis outcome,
- b) All the observations. Aggregate the data to get the arithmetic mean of Age per each kyphosis outcome,
- c) Observations that have Age greater than 50. Aggregate the data to get the count of observations per each kyphosis outcome,
- d) Observations that have Age smaller than 40. Aggregate the data to get the count of distinct Number values per each kyphosis outcome,
- e) Observations with Number smaller than 5. Aggregate the data to get the count of distinct Number values per each kyphosis outcome.

4. Write a SQL query on the iris dataset using the `squidf()` function to get the specified aggregation for:

- a) All the observations. Aggregate the data to get the sum of `Petal.Length` per each Species group,
- b) All the observations. Aggregate the data to get the arithmetic mean of `Petal.Width` per each Species group,
- c) Observations that have `Petal.Length` greater than 3.5. Aggregate the data to get the count of observations per each Species group,
- d) Observations that have `Sepal.Width` smaller than 3.4. Aggregate the data to get the count of distinct `Sepal.Length` values per each Species group,
- e) Observations with `Petal.Length` smaller than 1.8. Aggregate the data to get the count of distinct `Petal.Width` values per each Species group.

5. Create the following joins for the specific datasets:

- a) Inner join for `airlines` and `airports` from the `library(nycflights13)`
- b) Left join for `airlines` and `airports` from the `library(nycflights13)`
- c) Create the following datasets and perform a full outer join, left join, and right join:

Dataset A:

<b>customer_id</b>	<b>first_name</b>	<b>last_name</b>	<b>age</b>
<b>1001</b>	Paul	Smith	49
<b>1002</b>	John	Thomas	22
<b>1003</b>	Mary	Yip	15
<b>1004</b>	Thomas	Kurnicki	41
<b>1005</b>	Andrew	Rutky	32
<b>1006</b>	Jennifer	Yip	20

Dataset B:

<b>customer_id</b>	<b>transaction_type</b>	<b>total_value</b>
<b>1001</b>	1	55.99
<b>1001</b>	1	41.20
<b>1001</b>	2	69.01
<b>1003</b>	1	111.11
<b>1002</b>	1	19.00
<b>1002</b>	1	20.00
<b>1006</b>	2	28.40

d) Count how many observations are there in each of the new datasets created in a, b, and c.

6. Using the joins created in exercise 5.c calculate the following (use the `squidf()` function for your calculation) :

a) Median of `transaction_type` for all transactions above \$40,



- b) Median of transaction\_type for customers younger than 40,
- c) Sum of total\_value for all customers with truncation\_type equal to 2,
- d) Count distinct customer\_id,
- e) Sum of total\_value for each customer individually (group by customer),
- f\*\*) Frequency of transactions per customer,
- g\*\*) Frequency of transaction\_types per customer.

### **8. Visualizations and interactive plots**

#### 8.1 The basic plot() function

Even though GGPLOT2 and PLOTLY have become extremely popular over the past few years, R has some native plotting capabilities that are faster to use and implement. The most basic R function for plotting data is the plot() function (no need to call a library, the plot() is a native R function).

```
plot(x, y, type="selected_type")
```

The plot() function can plot a few different types of charts. The *selected\_type* can be:

- "p" for points,
- "l" for lines,
- "b" for both,
- "c" for the lines part alone of "b",
- "o" for both 'overplotted',
- "h" for 'histogram' like (or 'high-density') vertical lines.

The x from the function stands for the x-axis variable, whereas the y stands for the y-axis variable. These can be variables from a data frame and in this case they need to be specified in the data frame format `mydf$variable_name`.

## 8.1 The basic plot() function

Here's an example of a scatterplot using the plot() function:

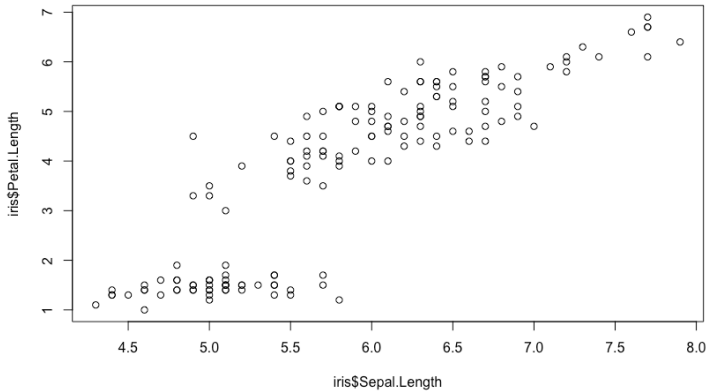


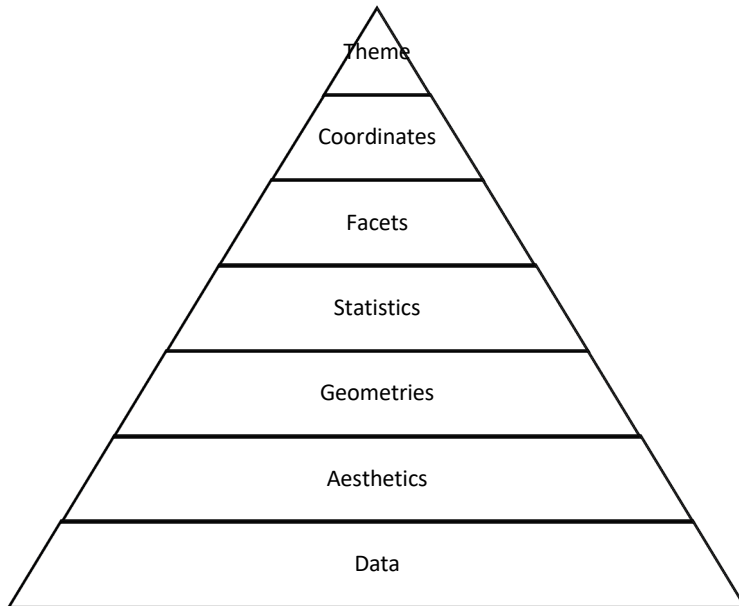
Chart 8.1.1 Scatter plot of the iris dataset [created by author].

The chart above is not very appealing. What is more, it has very limited options when it comes to adding layers. The most popular functions as an additional layer would be the lines() or points() functions.

```
lines(y, col="my_color")
```

### 8.2 GGPlot2 framework and library

GGPLOT2 provides more customizable plotting features than the regular `plot()` function. This plotting method is based on multiple layers of information that is passed to the plotting engine. The layers can be envisioned as parts of a pyramid, with the most important element at the base.



Picture 8.2.1 The GGPlot2 framework pyramid [created by author].

## 8.2 GGPLOT2 framework and library

The most important element of the ggplot framework, the **data**, points to the data frame that contains all the necessary data.

The **aesthetics** statement lists the X-axis and Y-axis variables and might specify the color variable if it doesn't depend on geometries.

**Geometries** specify the type of graphics for a given set of variables. Each geometry might have its own aesthetics. The most common geometries are:

- `geom_bar` – creates a bar chart,
- `geom_point` – creates a scatter plot,
- `geom_jitter` – creates a jitter plot,
- `geom_abline` – creates a line chart,
- `geom_text` – adds text for each coordinate specified in aesthetics,
- `geom_rect` – creates a plot with multiple rectangles,
- `geom_polygon` – creates a plot with polygons.

Multiple geometries can be combined on one graph. The “+” sign will combine the different layers:

```
ggplot(data=mydata, aes(x=x_var, y=y_var)) +  
  geom_point(aes()) +  
  geom_rug(aes())
```

When combining multiple geometries that have different statistics, the first `ggplot()` statement should be empty, allowing each consecutive `geom_x` statement to define new aesthetics:

```
ggplot() +  
  geom_point(data = df, aes(x = gp, y = y)) +  
  geom_point(data = ds, aes(x = gp, y = mean,  
    colour = 'red', size = 3)) +  
  geom_errorbar(data = ds, aes(x = gp, y = mean,  
    ymin = mean - sd, ymax = mean + sd),  
    colour = 'red', width = 0.4)
```

The three most important layers of this pyramid, described above can be found in any `ggplot` statement. The general `ggplot` template looks like this:

```
ggplot(df, aes(x = gp, y = y)) +  
  geom_point() +  
  geom_point(data = ds, aes(y = mean),  
    colour = 'red', size = 3)
```

The coordinates statement in the `ggplot` framework defines the location of the geometries. We have a few coordinates statements

## 8.2 GGLOT2 framework and library

that change the location of axes (source: <http://sape.inf.usi.ch/quick-reference/ggplot2/coord>, on 8/27/2018) :

- `coord_cartesian` - (default) cartesian coordinate system (x horizontal from left to right, y vertical from bottom to top),
- `coord_flip` - flipped cartesian coordinate system (x vertical from bottom to top, y horizontal from left to right),
- `coord_polar` - polar coordinate system; the x (or y) scale is mapped to the angle (theta),
- `coord_map` - various map projections.

```
ggplot(df, aes(x = gp, y = y)) +  
  coord_flip()
```

The **theme** statement will provide the “feel and look” for the charts introducing grid lines, background color, etc. GGLOT 2 has 2 built in themes:

- `theme_grey()` - the default theme, with a grey background,
- `theme_bw()` - white background.

The **facets** approach creates a matrix (grid) of plots. The main context in which the facets become useful is when there is

## 8.2 GGLOT2 framework and library

another variable or two that we want to group the data by. There are two main **facets**:

- `facet_wrap` – when there is one variable based on which we want to create a strip of charts and wrap it,
- `facet_grid` – when there are two variables that we want to use to create the matrix.

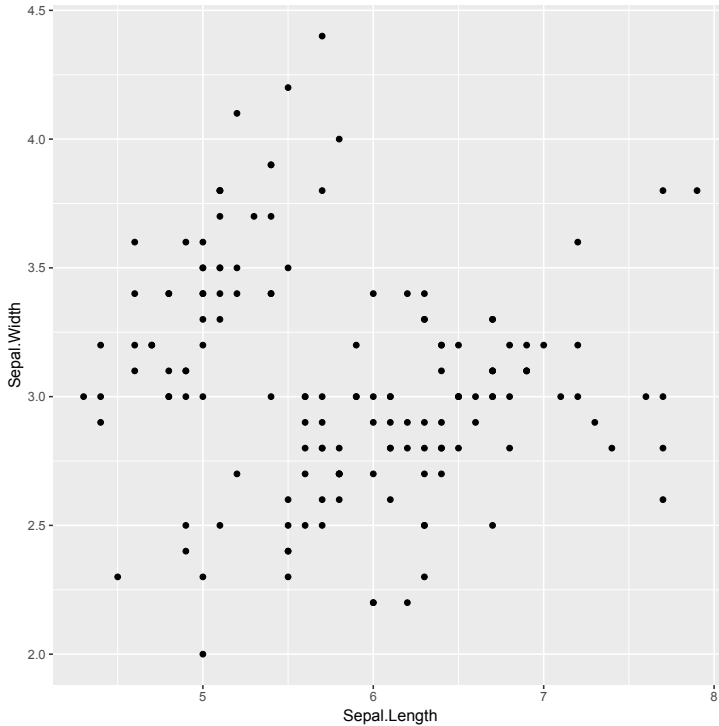
Additional GGLOT2 functions and tips can be found in the cheat sheet section in the appendix.

Let's consider the iris dataset that is available as a base dataset in R (no need to call any packages to load this data). In the below example, we want to create a scatter plot that shows the relationship between `Sepal.Width` and `Sepal.Length`. These two variables will be in our aesthetics layer. The `geom_point()` function will define the scatterplot type.

```
library(ggplot2)
ggplot(data=iris, aes(x=Sepal.Length, y=Sepal.Width))+
  geom_point()
```



## 8.2 GGPLOT2 framework and library



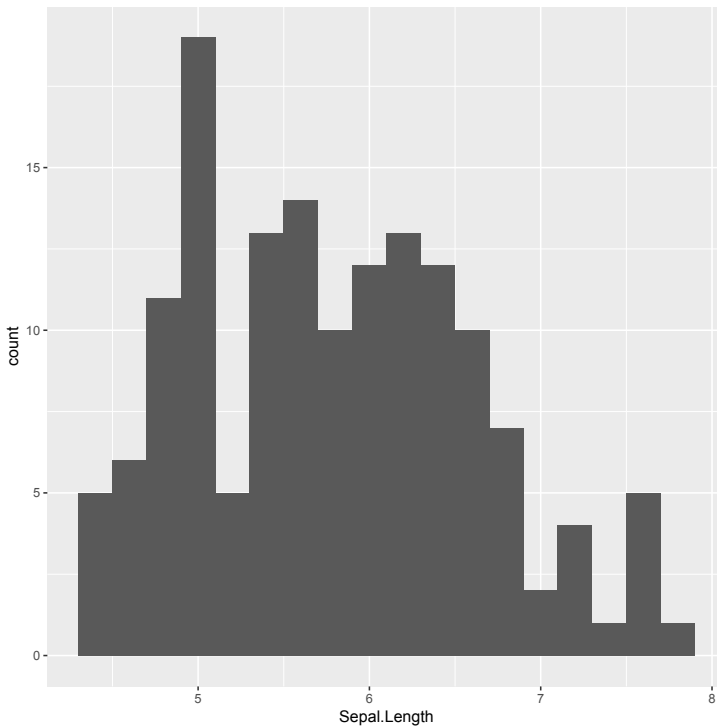
Picture 8.2.2 The GGPLOT2 scatterplot example with iris data [created by author].

The second example, showcased below, is a frequency histogram. Even though the chart looks like a bar chart, data scientists call it a frequency histogram because the height of the bar will indicate how many times a given value (on the X-axis) occurred in the particular variable. In this case, we use the `geom_histogram()` function to get a frequency histogram for the `Sepal.Length` variable.

## 8.2 GGLOT2 framework and library

```
library(ggplot2)
ggplot(data=iris, aes(Sepal.Length))+
  geom_histogram(binwidth=0.2)
```

The `geom_histogram()` function results in a chart show below.



Picture 8.2.3 The GGLOT2 frequency histogram example with iris data [created by author].

### 8.3 Interactive visuals with Plotly

Plotly is a next generation engine for creating plots in R. The most important advantage of using Plotly is its interactivity. Zooming in and out, pane selection, lasso selection are only a few features available in Plotly. What is more, Plotly is a great framework for creating interactive dashboards in Shiny. (Shiny is a package that allows creating dashboards and web based applications.)

First, Plotly has many similarities to the ggplot. It has a main `plotly()` statement followed by specific geometries. All these statements are combined using “%>%” instead of using the “+” sign from ggplot.

```
install.packages("plotly")  
library(plotly)
```

It is a good practice to create new objects with plotly charts and print them when finished. For a basic plotly box plot, the framework is as follows:

```
p <- plot_ly(data, x = ~x_var, color = ~y_var, type = "box")  
p
```

### 8.3 Interactive visuals with Plotly

Differently from the ggplot, plotly uses variable names with the approximation symbol “~”.

In order to change the type of the plot, the “type” option needs to be changed. The most basic chart types are as follows:

- scatter plots,
- box plots,
- bar plots.

In order to combine multiple layers of geometries and add traces or points, to the base plot, the “%>%” needs to be applied:

```
p<-plot_ly(data, x= ~x_var, color= ~y_var, type = "box")
p <- p %>% add_trace( x=x, y = y2 )
p <- p %>% layout(legend =
                    list( x = 0.5 ,
                          y=1,
                          bgcolor = '#F3F3F3' ))
p #this will print the p plot
```

Besides the “add\_trace” function that is presented above, there are other functions that help adding geometries to the chart:

- add\_lines – for adding lines to an existing chart,  
%>% add\_lines(y = ~fitted(loess(y\_var ~ x\_var))),  
line= list(color = 'rgba(7, 164, 181, 1)'),

```
name = "my_name")
```

- `add_markers` – to add custom elements, aka. “markers”,

```
%>% add_markers(y = ~y_var,
                 text = optional,
                 showlegend = FALSE)
```

- `add_ribbons` – to add a ribbon to a line chart.

```
%>% add_ribbons(data = my_data,
                 ymin = ~lower_band, ymax = ~upper_band,
                 line = list(color = 'rgba(7, 164, 181, 0.05)'),
                 fillcolor = 'rgba(7, 164, 181, 0.2)',
                 name = "my_name")
```

When it comes to customizing chart elements, the best function is called `marker()`. This function needs to be included in the main `plot_ly()` function.

```
plot_ly(data = my_data, x = ~x_var, y = ~y_var,
        marker = list(size = 10,
                      color = 'rgba(255, 182, 193, .9)',
                      line = list(color = 'rgba(152, 0, 0, .8)',
                                width = 2)))
```

When it comes to changing general chart display properties, the `layout()` function does all the work.

```
%>% layout(xaxis = list(title = 'x_name'),  
            yaxis = list(title = 'y_title'),  
            legend = list(x = 0.80, y = 0.90))
```

After running the plotly statements, the chart becomes available in the 4th, plot window. When the plot area is hovered over, a menu ribbon will pop-up in the upper right hand corner. A few icons will appear on the menu that do the following:

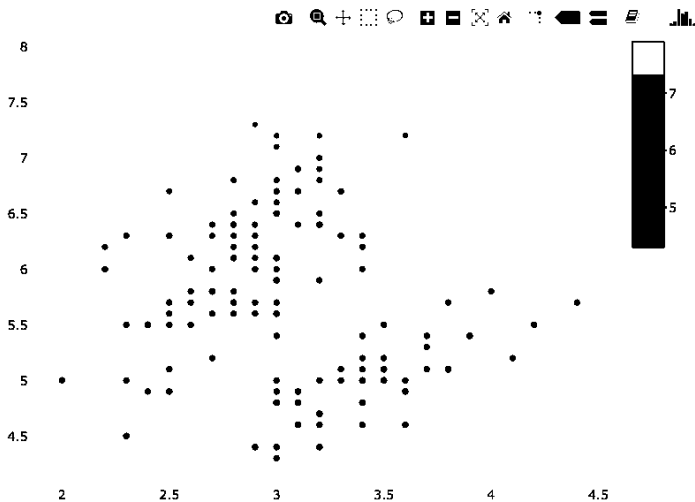
- zoom in and out,
- change the axes,
- provide a lasso selection,
- allow the user to add a tooltip when hovering over the chart elements and trace back the x and y values,
- return the original view.

In the example below, we are replicating the same chart that we created using `ggplot` in the previous chapter. It is a scatterplot that shows the relationship between the `Sepal.Width` and `Sepal.Length`.

```
library(plotly)
```

```
plot_ly(data=iris, x= iris$Sepal.Width,
        y= iris$Sepal.Length,
        color= iris$Sepal.Length,
        type = "scatter")
```

The above plotly code will produce a colorful scatter plot with all the plotly functionalities as described in this chapter such as zooming in and out, lasso selection, changing access, etc. Those options are visible in the upper right hand corner after hovering over the chart with your mouse.



Picture 8.3.1 The Plotly scatterplot example with iris data [created by author].

### 8.4 Exercises

1. Using the `plot()` function, create one scatter plot for each pair of variables and frequency histograms ( bar charts) for each variable separately:

- a) Sepal.Length and Petal.Width from the iris dataset,
- b) Sepal.Width and Petal.Length from the iris dataset,
- c) Number and Age from the kyphosis dataset,
- d) Age and Start from the kyphosis dataset,
- e) Age and total.value from the dataset created in exercise 5.c from chapter 7,
- f) Both variables in the cars dataset,
- g\*\*) The WFC stock and the SPY benchmark returns over the past 250 trading days (get the data using `quantmod` or `quandl`),
- h\*\*) The AAPL stock and the SPY benchmark returns over the past 250 trading days (get the data using `quantmod` or `quandl`).

2. Create ggplot plots for the following:

- a) Scatter plot of Sepal.Width and Petal.Width from the iris dataset,
- b) Matrix of scatter plots for all the variables in the iris dataset using `facets` layer,
- c) Scatter plot and a frequency histogram together on one plot for Age and Start from the kyphosis dataset,



d) Line chart for the price of AAPL stock from January 1<sup>st</sup> 2016 to September 1<sup>st</sup> 2018.

e\*\*) Log returns for any equity mutual fund for the period from January 1<sup>st</sup> 2015 to December 31<sup>st</sup> 2017.

3. Decide what chart type is best and create ggplots for the following data elements for the `titanic_train` dataset from the `library(titanic)`:

- a) Fare against age,
- b) Median age per pclass,
- c) Sum of fare per pclass and sex,
- d) Survived status against sex.

4. Recode the plots from exercise 2 in `plot_ly`. Analyze the plots using `plot_ly`'s features.

5. Recode the plots from exercise 3 in `plot_ly`.

a) Zoom in on any outliers and try to understand why these are outliers. Make screen shots of the zoomed in charts.

6. A hardware store chain has 30 stores. Each store generates daily revenue that is normally distributed over time (`rnorm()` function). Using a *for loop* to run a simulation of each store, plot the timeseries using the `plotly` framework.

## 8.4 Exercises

7. You are the data and analytics manager of a Fortune 500 company. Last week, you surveyed all the customer facing employees in the company's headquarters and noticed that their total compensation is based on two factors: years with the company and generated revenue. Fit a second degree polynomial function to predict the salary on the data.

salary	years_w_company	generated_rev
<b>216,344.77</b>	8	85,013.09
<b>242,226.87</b>	5	95,593.11
<b>29,487.10</b>	8	8,742.88
<b>216,246.57</b>	4	84,991.58
<b>20,905.48</b>	4	5,261.64
<b>181,123.01</b>	5.5	70,650.01
<b>129,937.99</b>	8	49,742.34
<b>293,054.30</b>	10	116,310.37
<b>39,758.87</b>	1.5	12,961.83
<b>111,967.53</b>	4.5	42,427.36
<b>10,990.51</b>	4.5	1,212.31
<b>33,739.11</b>	5.5	10,492.97
<b>200,791.05</b>	4	78,684.40
<b>273,329.40</b>	0.5	108,297.71
<b>111,807.44</b>	10	42,329.69
<b>109,325.17</b>	8	41,333.31
<b>258,285.53</b>	4	102,151.64
<b>81,401.51</b>	1	29,959.25
<b>258,794.62</b>	4	102,358.47

Plot the relationships between the variables side by side, using the plotly framework.

8. Using `ggplot()`, create frequency histograms (bar charts) for each variable from the Horse Colic Data Set that can be found at <https://archive.ics.uci.edu/ml/datasets.html>

9. Using `plotly()`, create timeseries plots to explain the evolvement of El Nino patterns over time. Consider humidity, temperature, subsurface temperatures. Are there any correlations that need to be explained (use scatter plots to understand correlations)? Use the El Nino Data Set from the <https://archive.ics.uci.edu/ml/datasets.html> website.

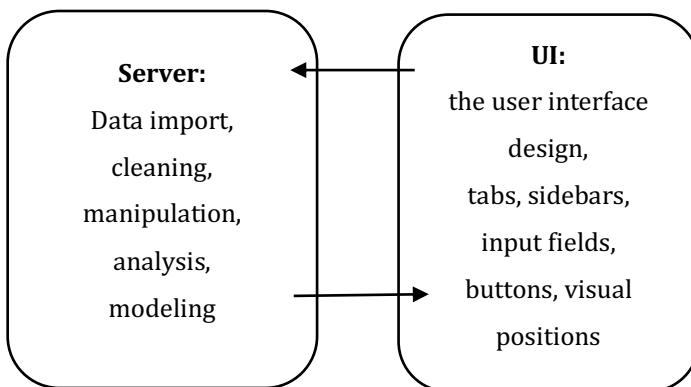
## **9. Building a dashboard in R Shiny**

### **9.1 Shiny as a framework**

Shiny is a great tool to build an interactive dashboard or web-based user interface for almost any visuals, tables, or text outputs. Shiny also allows the user to modify the data inputs and parameters.

The basic Shiny framework is based on two files:

- server file, that has all the data manipulation, and analysis logic,
- ui file, that stores the entire design of the user interface.



Picture 9.1.1 The general R Shiny framework [created by author].

To open a Shiny framework, click on the “new script” icon with the green plus sign in the upper left-hand side corner and click on the Shiny Web App option. This will open a pop-up window with a few options. The first one will ask the user to provide an application name. The application name is not very important at this point and we can provide any random name. The second selection is Single File or Multiple File. It is recommended to work in the “Multiple File” mode. The third option needs the user to specify the directory where the files will be saved. Both files, the Server.R and UI.R need to be saved in the same directory.

Two new scripts will open in the “script window” after clicking OK in the Shiny pop up window. The two new files are called:

- Server.R,
- UI.R.

These two files will have default templates that can be easily modified for any custom project.

### 9.2 Server and UI interaction

Even though the Server.R and UI.R are separate files, they are required to have common elements.

The server can do all the work related to importing and manipulating/transforming the data. However, in the end it has to create objects that are readable by the UI.R file. These objects can be created using the following functions:

- `renderImage()` – creates a static image for the UI,
- `renderPlot()` – takes simple plots and ggplots,
- `renderPrint()` – builds an object with text or console outputs,
- `renderTable()` – stores a data frame,
- `renderText()` – stores static text,
- `renderPlotly()` – stores plotly charts.

These functions work just like the user defined functions and need to return an object that has the specific type.

What is more, each of these objects has to be named using the following template:

```
output$name <- renderTable()
```

Each server function listed above has a matching function for the UI.R file:

- `renderImage()` – `imageOutput()`,
- `renderPlot()` – `plotOutput()`,
- `renderPrint()` – `verbatimTextOutput()`,
- `renderTable()` – `tableOutput()`,

- `renderText()` – `textOutput()`,
- `renderPlotly()` – `plotlyOutput()`.

The UI.R file will use the *name* defined in the Server.R as function input for the corresponding UI function.

### 9.3 Basic server functions

There are many server side functions that give outputs for the UI. It would be impossible to list out all of the functions in this chapter. However, we'll take a look at the `renderPlot()` function to understand how Shiny creates objects. Please remember that each render function needs to be saved as an **OUTPUT** object using the following template:

```
output$obj_name <- renderPlot()
```

The `renderPlot()` function has a very detailed description with many options that can be found on the shiny website [source: [www.shiny.rstudio.com](http://www.shiny.rstudio.com)]. Despite all the available options, using the function below is more than enough.

```
output$myplot <- renderPlot(expr, width = "auto", >  
                           height = "auto")
```

Where:

width – is the width of the plot, can be “auto” or a fixed dimension,  
height – is the height of the plot, and can be fixed as well,  
expr- are all the regular R expressions that create the plot, this could be a plot(x) function or a set of functions that lead to creating a ggplot object. In the case of using a ggplot framework, the object needs to be saved in the global environment (e.g. P ) and then called at the end of the code, like this:

```
output$myplot <- renderPlot(  
  P <- ggplot(data=xyz, aes(x=x, y=y))+  
    geom_lines(...)  
  print(P)  
, width = "auto", height = "auto")# closing renderPlot
```

### 9.4 UI elements and user inputs

The UI.R will have two types of objects/functions:

- user inputs – multiple choice boxes, date inputs, text inputs, radioButtons, actionButtons, etc,
- server output – any function like plotOutput(), textOutput(), etc., that refers to an object created in the server.



The “server output” functions have been listed at the end of chapter 9.2. These `xyzOutput()` functions require an object name that was created in the server using a corresponding “server side” function.

On the other hand, the user inputs will create objects that the server can easily read using a template:

```
input$ui_input_name
```

The most basic user inputs include:

- `actionButton()`,
- `checkboxGroupInput()`,
- `checkboxInput()`,
- `dateInput()`,
- `fileInput()`,
- `radioButtons()`,
- `textInput()`,
- `sliderInput()`.

More functions for the UI.R and Server.R can be found in the Appendix in the R Shiny Cheat Sheet.

One of the most important features that R Shiny has is the ability to modify the layout of the UI. In order to manage the end user

experience, Shiny allows to create tabs, menus, sidebars, etc. , just like in HTML.

There are 5 basic ways to organize panels:

- sidebarLayout – the most popular

```
ui <- fluidPage(  
  sidebarLayout(  
    sidebarPanel(),  
    mainPanel()  
  )#closing sidebarLayout  
)#closing fluidPage
```

- splitLayout

```
ui <- fluidPage(  
  splitLayout(  
    object1, object2,.....  
  )#closing splitLayout  
)#closing fluidPage
```

- verticalLayout

```
ui <- fluidPage(  
  verticalLayout(  
    object1, object2,.....  
  )#closing verticalLayout  
)#closing fluidPage
```

- `flowLayout`

```
ui <- fluidPage(  
  flowLayout(  
    object1, object2,.....  
  )#closing flowLayout  
)#closing fluidPage
```

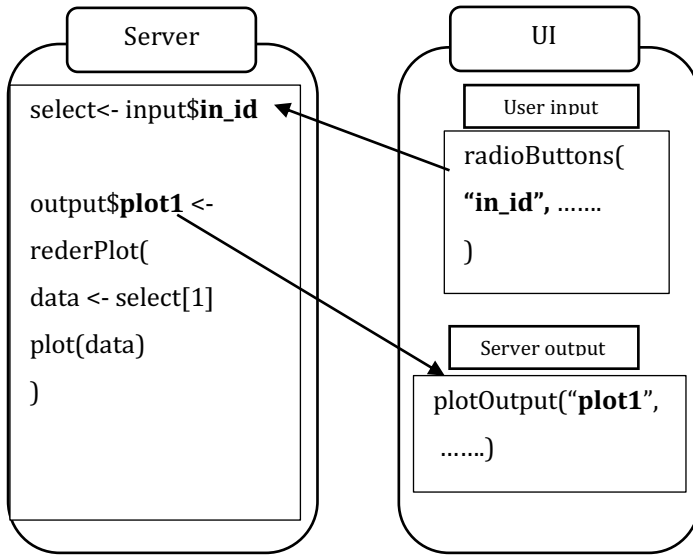
- `fluidRow`

```
ui <- fluidPage(  
  fluidRow(column(), column()),  
  fluidRow(column(),column())  
)#closing fluidPage
```

What is more, each `Layout` can contain multiple tabs. The `tabsetPanel` function controls the dynamics of tabs.

```
ui <- fluidPage(tabsetPanel(  
  tabPanel("tab1", content),  
  tabPanel("tab2", content)  
)#closing tabsetPanel  
)#closing fluidPage
```

Finally, the graph below visualizes how the objects and information flow from the UI to server and back to the UI.



Picture 9.4.1 Information flow between UI and SERVER [created by author].

### 9.5 Exercises – R Shiny project

Project 1. Create a web-based application that uses the kyphosis dataset and allows the user to select two variables (create two drop down selection boxes). The application creates a ggplot scatterplot for any given set of two variables.

Project 2. Using the titanic\_train dataset from the library(titanic), create a web based application with a sidebar and two tabs. The sidebar has:

## 9.5 Exercises – R Shiny project

- a) a `sliderInput()` that allows to subset the age variable,
- b) `radiobuttons` that subset the sex variable

The first tab will have a `plotly()` graph of Age and Fare for all remaining observations, based on the selections from the sidebar. Tab two will have a `logistic regression summary()` for a model that predicts if the person survived.

Project 3\*\*. Recode project 2 using `library(shinydashboard)` approach. Make sure to use two or three `valueBox()` objects on the landing page and provide some basic statistics for the titanic dataset. Use appropriate icons in all `valueBox()` objects.

## 10.1 R Markdown setup and basic features

### **10. Report automation in R Markdown**

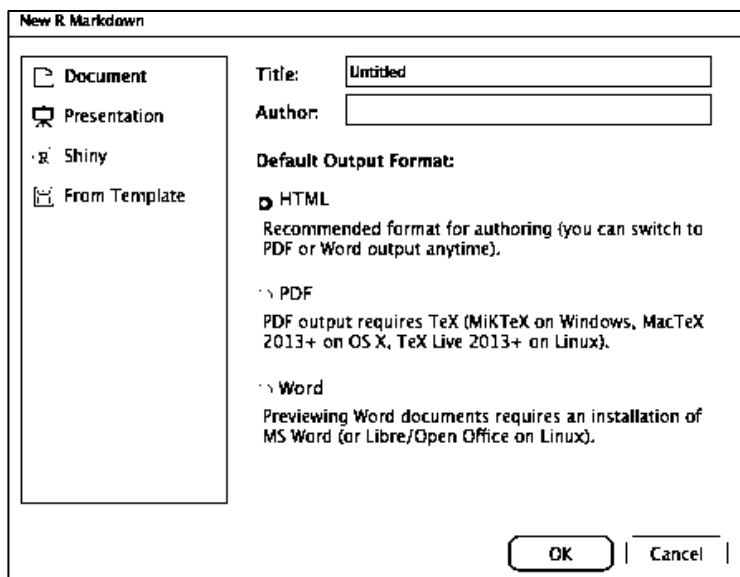
#### 10.1 R Markdown setup and basic features

R Markdown is a great tool for report automation. It can generate reports in a few formats:

- HTML - same styling capabilities as regular HTML websites and CSS styling can be applied as well,
- PDF – less styling options, easier to distribute via email,
- other – including a MS Word document (these need additional packages and will not be covered in this chapter).

A new Markdown project can be opened by clicking the “new script” icon with the green plus sign in the upper left corner. When the drop down list comes up, select the “R Markdown...” option. A pop-up window will appear with a few fields to fill out. In this chapter we’ll use the “Document” option from the left-hand side list.

## 10.1 R Markdown setup and basic features



Picture 10.1.1 R Markdown Wizard [created by author].

We highly recommend that you explore the other options that are listed, such as “Presentation” or “From Template”. In the other fields, you will need to provide the document name or title, author, and the document type using the radio buttons. Select the HTML document type for your first R Markdown project.

## 10.2 Code chunks and types in Markdown

### 10.2 Code chunks and types in Markdown

When all the steps from chapter 10.1 are completed, a script template with the following code chunks will open:

```
---
title: "Untitled"
output: html_document
---

```${r setup, include=FALSE}
knitr::opts_chunk$set(echo = TRUE)
```

## R Markdown

This is an R Markdown document. Markdown is a simple
formatting syntax for authoring HTML, PDF, and MS
Word documents. For more details on using R Markdown
see <http://rmarkdown.rstudio.com>.

When you click the Knit button a document will be
generated that includes both content as well as the
output of any embedded R code chunks within the
document. You can embed an R code chunk like this:
```



## Markdown

```
```{r cars}
summary(cars)
```
```

The first few lines, separated from the other code using --- (dashes), with the title and output statements are called the header. The header contains basic information about the Markdown document such as the title, author name, date, and type (HTML, PDF, or other). It might also have a reference line to a .css file with all the formatting styles for a HTML document.

The second chunk is where the R code is stored. This chunk will execute just like any regular R code and output only the object that is printed. The printed output can be a plot, model summary, data frame, etc. This chunk has to be surrounded by ``` and the {r *options*} indicates if the errors and warnings need to be printed.

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(echo = TRUE)
```
```

## 10.3 Exercises – R Shiny project \*\*

In the initial Markdown script, the following statement is included:

```
## R Markdown
```

Markdown understands #, ##, ### as styling references and will print the following text as a header in the final output document.

Finally, Markdown can print simple text to the final output document. It doesn't need any special characters around it and will print the text using default font.

Once your Markdown script is ready and has all the code chunks, click the drop down next to the “Knit” icon and select the appropriate document type. This will produce the final output document.

### 10.3 Exercises – R Shiny project \*\*

Project 1. Using the scatterplots from Project 1 in chapter 9. Create a report with scatterplots for all pairs of variables and provide short descriptions for each scatter plot.

Project 2. Use the plot and model summary from Project 2 in chapter 9. Create a report that will describe the plot and the model summary. Focus your descriptions on the significance of variables.

## **11. List of most useful functions**

### **11.1 Mathematical**

log(x),logb(),log10(),log2(),exp(),expm1(),log1p(),sqrt()  
 cos(),sin(),tan(),acos(),asin(),atan(),atan2()  
 cosh(),sinh(),tanh(),acosh(),asinh(),atanh()  
 union(),intersect(),setdiff(),setequal()

#### **Basic operators:**

+, -, \*, /, ^, %%, %/%

#### **Comparison operators:**

<, >, <=, >=, ==, !=

#### **Principal component analysis, eigenvectors and eigenvalues:**

eigen(), princomp()

#### **Calculus:**

deriv()

integrate()

sqrt(), sum()

### **11.2 Data manipulation**

read\_csv()

`read_tsv()`  
`read_delim()`  
`read_fwf()`  
`read_table()`  
`as.data.frame()`  
`as.numeric()`  
`as.character()`  
`as.Date()`  
`na.omit()` -excludes all the rows with at least one missing  
`is.na()` -tests an object for being empty  
`t()` -transposes a matrix  
`gsub()` -replaces a given string with a new string  
`subset()` -subsets a data frame based on a value  
`which()` -gives indexes based on a condition  
`extract()` -Extract one column into multiple columns.  
`gather()` -Gather columns into key-value pairs.  
`nest()` -nest a few columns together  
`separate()` -Separate one column into multiple columns.  
`separate_rows()` -Separate a collapsed column into multiple rows.  
`unite()` -Unite multiple columns into one.  
`drop_na()` -Drop rows containing missing values  
`fill()` -Fill in missing values.  
`replace_na()` -Replace missing values

## 11.3 Regression, optimization, and fitting

### 11.3 Regression, optimization, and fitting

|            |  |
|------------|--|
| lm()       | -linear modeling                       |
| glm()      | -generalizes linear modeling           |
| nls()      | -non linear leased squares             |
| predict()  | -gives prediction for most models      |
| summary()  | -provides a summary for a model        |
|            |  |
| optim()    | -general optimization                  |
| optimize() | -1 dimentional optimization            |
| nlm()      | -non linear minimization               |
| nlmLM()    | -non linear minimization, more options |

### 11.4 Descriptive statistics

mean()  
sum()  
median()  
sd()  
colMeans()  
percentile()  
summary()

### **12. Data mining with MongoDB (noSQL)**

In one of the previous chapters, we've explained how to use SQL to combine, analyze, and subset data. Those SQL queries work very well on relational data that is linked by primary and secondary keys. However, more and more data is stored in non-relational databases, that use the NoSQL language. Non-relational databases can be created to host different objects such as documents, keys, graphs, pictures, etc. and do not require defining relationships between the data objects.

#### **12.1 Installing MongoDB**

MongoDB is a database that combines both concepts, the capabilities of relational databases, and the flexibility of NoSQL. What is more, Mongo DB is a free, open source database that runs on a NoSQL engine. The data behind it is saved in a JSON format. The first step to start using MongoDB is to install the software. The drivers can be installed on a macOS, Windows, or Linux and come in two editions:

- Community Edition – is a free open source version,
- Enterprise – provides more security and encryption but has a cost associated with it.

## 12.1 Installing MongoDB

There are different steps for installing MongoDB on different operating systems. A full list of instructions for each operating system can be found at <https://docs.mongodb.com/manual/installation/>

For the macOS, the instructions are as follows [source: <https://docs.mongodb.com/manual/tutorial/install-mongodb-on-os-x/>, on 8/24/2018] :

Install MongoDB Community Edition Manual:

### Step 1:

Download the MongoDB .tar.gz tarball.

- Download the tarball for your system from the MongoDB Download Center  
<https://www.mongodb.com/download-center#production>

### Step 2:

Extract the files from the downloaded archive.

- Submit the following statement in your terminal:

```
tar -zxvf mongodb-osx-ssl-x86_64-4.0.1.tgz
```

## 12.2 Installing MongoDB Community

with Homebrew

### Step 3:

Ensure the binaries are in a directory listed in your PATH environment variable. The MongoDB binaries are in the **bin/** directory of the tarball. You must either:

- Copy these binaries into a directory listed in your PATH variable such as **/usr/local/bin**,
- Create symbolic links to each of these binaries from a directory listed in your PATH variable, or
- Modify your user's PATH environment variable to include this directory.

For example, you can add the following line to your shell's initialization script (e.g. **~/.bashrc**):

```
Export PATH=<mongodb-install-directory>/bin:$PATH
```

Replace **<mongodb-install-directory>** with the path to the extracted MongoDB archive.

### 12.2 Installing MongoDB Community with Homebrew

The instructions below are copied from the official MongoDB website: <https://docs.mongodb.com/manual/tutorial/install-mongodb-on-os-x/>



## 12.2 Installing MongoDB Community

with Homebrew

Homebrew installs binary packages based on published “formulae.” This section describes how to update brew to the latest packages and install MongoDB Community Edition. Homebrew requires some initial setup and configuration, which is beyond the scope of this document.

### Step 1:

Update Homebrew’s package database.

In a system shell, issue the following command:

```
brew update
```

### Step 2:

Install MongoDB.

You can install MongoDB via brew with several different options.

Use one of the following operations:

- Install the MongoDB Binaries  
To install the MongoDB binaries, issue the following command in a system shell:

```
brew install mongodb
```

- Install the Latest Development Release of MongoDB

## 12.3 Running MongoDB

To install the latest development release for use in testing and development, issue the following command in a system shell:

```
brew install mongodb --devel
```

### 12.3 Running MongoDB

The following steps will boot Mongo DB:

#### **Step 1:**

Create the data directory.

Before you start MongoDB for the first time, create the directory to which the mongod process will write data. By default, the mongod process uses the `/data/db` directory. If you create a directory other than this one, you must specify that directory in the `dbpath` option when starting the mongod process later in this procedure.

The following example command creates the default **`/data/db` directory**:

```
mkdir -p /data/db
```

### Step 2:

Set permissions for the data directory.

Before running `mongod` for the first time, ensure that the user account running `mongod` has read and write permissions for the directory.

### Step 3:

Run MongoDB.

To run MongoDB, run the **mongod** process at the system prompt. If necessary, specify the path of the **mongod** or the data directory. See the following examples:

- Run without specifying paths

If your system `PATH` variable includes the location of the `mongod` binary and if you use the default data directory (i.e., `/data/db`), simply enter `mongod` at the system prompt:

```
| mongod
```

- Specify the path of the `mongod`

If your `PATH` does not include the location of the `mongod` binary, enter the full path to the `mongod` binary at the system prompt:

```
| <path to binary>/mongod
```

- Specify the path of the data directory  
If you do not use the default data directory (i.e., /data/db), specify the path to the data directory using the --dbpath option:

```
mongod --dbpath <path to data directory>
```

### Step 4:

Verify that MongoDB has started successfully.

Verify that MongoDB has started successfully by checking the process output for the following line:

```
[initandlisten] waiting for connections on port 27017
```

The output should be visible in the terminal or shell window.

You may see non-critical warnings in the process output. As long as you see the log line shown above, you can safely ignore these warnings during your initial evaluation of MongoDB.

### Step 5:

Begin using MongoDB.

Start a mongo shell on the same host machine as the mongod. Use the --host command line option to specify the localhost address (in this case 127.0.0.1) and port that the mongod listens on:

```
mongo --host 127.0.0.1:27017
```

Later, to stop MongoDB, press Control+C in the terminal where the mongod instance is running.

### 12.4 Using MongoDB from R

The best package for establishing a connection with MongoDB from R is the mongolite library.

```
install.packages("mongolite")  
library(mongolite)
```

If the MongoDB that you're trying to access is on your computer, you must start it by running the following in your shell:

```
mogod
```

Once we have our database up and running we will have to create our first connection, database, and collection. The `mongo()` function will create a collection:

```
my_collection = mongo(collection =  
                        "collection_name",
```

```
db = "db_name")
```

## 12.5 Uploading files to MongoDB

The easiest way to upload data to a collection defined in chapter 12.4 is to use the `$insert` options.

```
my_collection$insert(my_data)
```

## 12.6 Managing non-relational data using NoSQL

MongoDB uses JSON based syntax to manage the data. There are a few basic functions, including the **\$find** function, that take JSON syntax as arguments.

```
my_new_data <- my_collection$find(  
  query = '{"field1" : "field1_val", "field3" : { "$lt" : 100 } }',  
  fields = '{"field1" : true, "field2" : true}',  
  limit = 5  
)
```

The above query filters (**\$find**) *field1* based on the *field1\_val*. The “query” statement has all the filters, whereas the “fields” section lists all variables that are supposed to be printed. The “limit” puts

a limit on the number of rows that get carried over to the new table.

When it comes to aggregating data (**\$aggregate**), the JSON syntax requires one of the following aggregation names such as **\$sum**, **\$avg**, **\$count**, etc. and the **\$group** reference.

```
my_new_data <- my_collection$aggregate(  
  '[{"$group":{"_id":"$carrier","count":{"$sum":1},  
    "average":{"$avg":"$distance"}}}]',  
  options = '{"allowDiskUse":true}'  
)
```

MongoDB allows to use a more flexible aggregation method called “mapreduce”. Mapreduce uses JavaScript language to call more complex aggregations in MongoDB. Unless you know JavaScript, it is recommended to pull raw data from MongoDB to R and perform operations in R using R functions.

### 12.7 Exercises

1. Move the following datasets to the MongoDB and create new collections:

- a) iris,
- b) kyphosis.

2 Get the following from the kyphosis collection from MongoDB:

- a) All the variables and all the observations,
- b) Just the Number and Age variables and all the observations,
- c) All the variables and only the observations that have Age greater than 50,
- d) Just the Number and Age variables and only the observations that have Age smaller than 40,
- e) Just the Kyphosis variable for all the observations with Number smaller than 5.

3. Get the following from the iris collection from MongoDB:

- a) All the variables and all the observations,
- b) Just the Sepal.Length and Sepal.Width variables and all the observations ,
- c) All the variables and only the observations that have Sepal.Length greater than 4,
- d) Just the Pedal.Length and Sepal.Width variables and only the observations that have Sepal.Length smaller than 4,
- e) Just the Sepal.Length variable for all the observations with Sepal.Width smaller than 4.

4. Get the following from the kyphosis collection from MongoDB:

- a) All the observations. Aggregate the data to get the sum of Age per each kyphosis outcome,



- b) All the observations. Aggregate the data to get the arithmetic mean of Age per each kyphosis outcome,
- c) Observations that have Age greater than 50. Aggregate the data to get the count of observations per each kyphosis outcome,
- d) Observations that have Age smaller than 40. Aggregate the data to get the count of distinct Number values per each kyphosis outcome,
- e) Observations with Number smaller than 5. Aggregate the data to get the count of distinct Number values per each kyphosis outcome.

5\*\*. Get the following from the iris collection from MongoDB:

- a) All the observations. Aggregate the data to get the sum of Petal.Length per each Species group,
- b) All the observations. Aggregate the data to get the arithmetic mean of Petal.Width per each Species group,
- c) Observations that have Petal.Length greater than 3.5. Aggregate the data to get the count of observations per each Species group,
- d) Observations that have Sepal.Width smaller than 3.4. Aggregate the data to get the count of distinct Sepal.Length values per each Species group,
- e) Observations with Petal.Length smaller than 1.8. Aggregate the data to get the count of distinct Petal.Width values per each Species group.

## 13.1 Importing data and creating a text corpus

### **13. Text analytics**

#### 13.1 Importing data and creating a text corpus

Text analytics requires uploading multiple data files, such as .pdf or .doc files. Each file is treated as a separate observation.

To upload multiple .pdf files we will use the library(tm):

```
install.packages("tm")  
library(tm)
```

Move all the .pdf files to one location that is somewhere locally, on your computer and run the following function to get all the .pdf file names in the specified directory.

```
files <- list.files(path="my_path",pattern = "pdf$")
```

Once you have created a list with all the document names, you'll need to import all the documents into the R environment. To import the .pdf files, we'll need to create an "engine" function that can be used when creating a text corpus.

```
install.packages("pdftools")  
library(pdftools)
```

## 13.1 Importing data and creating a text

corpus

```
Rpdf <- readPDF(control = list(text = "-layout"))
```

The `readPDF()` function listed above creates a user defined function, the “engine”, and saves it in the environment. This function does not import the data.

Next, we’ll use the “engine” function to import the Corpus (a Corpus is an object with all the data from all the documents):

```
my_corpus <- Corpus(URISource(files),  
  readerControl = list(reader = Rpdf))
```

The process for importing text files is slightly different from the .pdf importing process explained above. Text files are much easier to read as they store information about each character and string that is included in the file.

The main difference between importing .pdf and .doc file is the reader options. For .pdf files, we have to create a `readPDF()` function, whereas the `readDOC` reader is used for .doc files.

```
files1 <- list.files(path="my_path", pattern = ".doc")  
ex_eng <- Corpus(URISource(files1),  
  readerControl=list(reader=readDOC,  
    language='en_CA',  
    load=TRUE))
```

### 13.2 Term-document matrix (TDM)

In chapter 13.1, we've explained how to create a Corpus using multiple .pdf and .doc files. With such a data structure we should be able to convert all the documents from the Corpus into a **term-document matrix (TDM)**. The TDM is a large matrix with all the tokens (term "token" stands for one "word") found in all documents with frequencies. We can use the following function from the library(tm):

```
my_tdm <- TermDocumentMatrix(my_corpus,  
  control = list(removePunctuation = TRUE,  
                 stopwords = TRUE,  
                 tolower = TRUE,  
                 stemming = TRUE,  
                 removeNumbers = TRUE))
```

The TermDocumentMatrix() has many options that allow to clean up the Corpus. It is highly recommended to build the TDM with options that exclude all punctuation, upper/lower case variations of the same word, and numbers.

Alternatively, the TDM can be cleaned up using the tm\_map() function:

## 13.2 Term-document matrix (TDM)

```
my_tdm <- tm_map(my_tdm, stripWhitespace)
my_tdm <- tm_map(my_tdm,
content_transformer(tolower))
my_tdm <- tm_map(my_tdm, removeWords,
stopwords("english"))
```

The `tm_map()` with `removeWords`, `stopwords()` option will remove all the words that do not make sense to analyze, such as: I, we, he, she, is, are, being, etc.

The `inspect()` function will analyze the TDM. What is more, this function can be subset the TDM just like any other matrix using indexes for rows and columns, or column names (column names or row names will have the tokens). The outcome will print a frequency table (how many times a word was used) for all the tokens.

```
inspect(my_tdm)
```

Other useful functions used to analyze the TDM are:

```
findFreqTerms(my_tdm, frequency_count)
findAssocs(my_tdm, "my_word", correlation)
```

### 13.3 Get sentiments from tidytext

The first function will print all the words for a given frequency.

The *frequency\_count* has to be an integer greater than 0.

The second function will print all the words that are associated with *my\_word* at a given *correlation* (number between 0 and 1).

#### 13.3 Get sentiments from tidytext

The very well-known package called tidytext offers sentimental token classification. It will assign a negative or positive label to all the tokens that are available in the library.

First, the TDM had to be transformed using the tidy() function:

```
library(dplyr)
library(tidytext)

tidy_tdm <- tidy(my_tdm)
```

Next, we need to left join the sentiment data from library(tidytext) using the get\_sentiments() function. The following code joins the two:

```
tdm_sentiments <- tidy_tdm %>%
  inner_join(get_sentiments("bing"), by = c(term =
"word"))
```

### 13.3 Get sentiments from tidytext

After joining the sentiments to the TDM data, we can run basic analytics to find the most negative documents in our TDM[source: [https://cran.r-project.org/web/packages/tidytext/vignettes/tidying\\_casting.html](https://cran.r-project.org/web/packages/tidytext/vignettes/tidying_casting.html) on 8/10/2018]:

```
library(tidyr)

tdm_sentiments %>%
  count(document, sentiment, wt = count) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative) %>%
  arrange(sentiment)
```

We can also create a visual that gives frequencies of the most positive and most negative tokens in our TDM [source: [https://cran.r-project.org/web/packages/tidytext/vignettes/tidying\\_casting.html](https://cran.r-project.org/web/packages/tidytext/vignettes/tidying_casting.html) on 8/10/2018]:

```
library(ggplot2)

tdm_sentiments %>%
  count(sentiment, term, wt = count) %>%
```

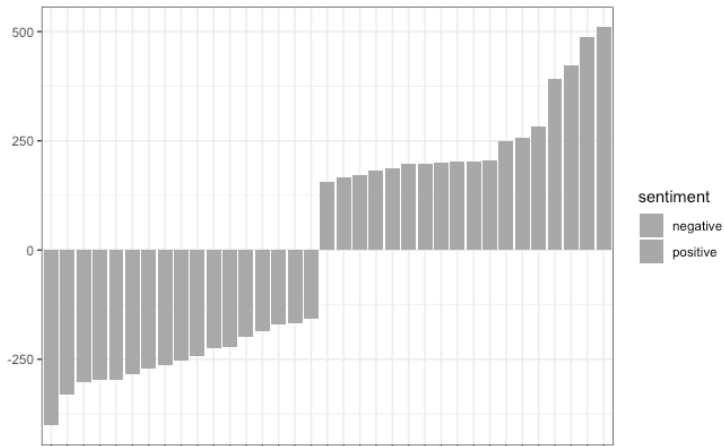
```

filter(n >= 150) %>%
mutate(n = ifelse(sentiment == "negative", -n, n)) %>%
mutate(term = reorder(term, n)) %>%

ggplot(aes(term, n, fill = sentiment)) +
geom_bar(stat = "identity") +
theme(axis.text.x = element_text(angle = 90, hjust = 1))

```

The ggplot statement will give the following bar chart:



Picture 13.3.1 Negative and positive token frequency [created by author].

### 13.4 Creating word clouds

Word clouds are probably the worst type of plots used to visualize text frequencies. Nevertheless they are used very often and need to be described in this chapter.



Since word clouds are so popular, a package called `wordcloud` was created. The package has a function called `wordcloud()` that creates the word cloud based on the tokens and the frequencies. Before we introduce the `wordcloud()` function, we need to manipulate the TDM so that it becomes a sorted data frame. Here are the following steps:

- Convert the TDM to a matrix,
- Sort the matrix by decreasing frequency of the tokens,
- Create a data frame with tokens and frequencies.

The following template can be used to manipulate the TDM object (we'll use the `my_tdm` object created in chapter 13.3):

```
p <- as.matrix(my_tdm)
v <- sort(rowSums(p),decreasing=TRUE)
final_df <- data.frame(word = names(v),freq=v)
```

The final step is to use the transformed data frame called `final_df` in the `wordcloud()` function.

```
wordcloud(words = final_df $word,
          freq = final_df $freq, min.freq = 1,
          max.words=200, random.order=FALSE,
          rot.per=0.35, colors=brewer.pal(8, "Dark2"))
```

### 13.5 Bayesian text classification model

Some documents might be classified into groups. In most cases the groups can be labeled as 1 for all the “yes”, “present”, etc. groups and 0 for the “no”, “absent” groups. In this case we can build a supervised learning model based on Bayesian inference.

First, we need to make sure that our TDM is a matrix and that we have a “y” (label) variable for each document in this matrix.

```
tdm_matrix <- as.matrix(my_tdm)
tdm_matrix1 <- cbind(tdm_matrix, c(0,1))
colnames(tdm_matrix1)[ncol(tdm_matrix1)] <- y
tdm_matrix1 <- as.data.frame(tdm_matrix1)
tdm_matrix1$y <- as.factor(tdm_matrix1$y)
```

The `tdm_matrix1` is a data frame with binary labels, called the “y” variable. The “y” variable needs to be a factor to work in the next step. In the next step, we will fit a Bayesian model on the data to get a predictive framework.

```
Library(caret)
baysian1 <- train(y ~ ., data = tdm_matrix1,
                  method = 'bayesglm')
```

Once we fit the model, we can use the `predict()` function to predict the “y” label for any new documents.

### 13.6 Exercise

1. Using the AssociatedPress documents from `library(topicmodels)`, analyze the TDM and look for most frequent tokens. Hint: after installing the `topicmodels` package, run the following code:

```
data("AssociatedPress", package = "topicmodels")
```

2. Using the TDM created in exercise 1, find the sentiments using `library(tidytext)`.

- a) Plot a word cloud for all the tokens that have a positive sentiment,
- b) Plot a word cloud for all the tokens that have a negative sentiment.

3. Create a Bayesian text classification model based on the documents used in exercise 2.

- a) What are the model parameters,
- b) What can you tell about the model significance.

## **14. Appendix**

rstudio.com]

## 14.1 READR cheat sheet [source: rstudio.com]

## Data Import : : CHEAT SHEET

**RStudio** is built around **big data** stored in **tables**, which is a structured data format. The first step of this sheet shows how to read data files into R with **readr**.

The **readr** side shows how to create tables with **tbls** and to import **tbls** data with **tblr**.

## OTHER TYPES OF DATA

Try one of the following packages to import other types of files: **Shape** and **SAS files**

- **readr** - read files (xls and xlsx)
- **DBI** - databases
- **jsonlite** - json
- **xmllite** - XML
- **httr** - Web APIs
- **readr** - **HTML** (Web scraping)

## Save Data

Save a **tbl** object to **path**, a file path, as:

```

# Create a dummy file
write_csv(tbl, "tbl.csv")

# Write with arbitrary delimiter
write_delim(tbl, delim=";", na="NA",
  append=FALSE, col_names=TRUE)

# CSV for excel
write_excel_csv(tbl, path="tbl.csv", append=
  FALSE, col_names=TRUE)

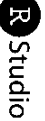
# Write to file
write_file(tbl, path="tbl.csv", append=FALSE)

# Write vector to file, one element per line
write_lines(tbl, path="tbl.csv", append=FALSE,
  delim=";", na="NA", progress=FALSE)

# Object to RDS file
write_rds(tbl, path="tbl.rds", na="NA",
  delim=";", na="NA", progress=FALSE)

# Tab delimited file
write_tsv(tbl, path="tbl.tsv", na="NA",
  delim=";", na="NA", progress=FALSE)

```



## Read Tabular Data

These functions share the common arguments:

```

read_rdf, col_types = TRUE, col_types = NULL, locale = default, na = c("NA",
  "NA"), progress = interactive(),
  n_max = 1000, progress = interactive()

```

6.0k

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

1.2.3

4.8 NA

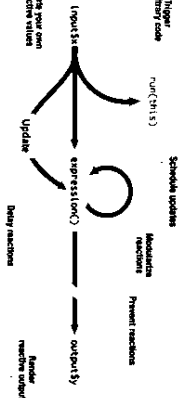








Reactive values work together with reactive functions. Call a reactive value from within the arguments of one of these functions to avoid the error operation not allowed without an active reactive context.



# example snippets  
 ui <- FluidPage({  
 textInput("m1", "m")  
 reactiveValues({  
 (see front page)

```

1 public static void main(String[] args) {
2     ReactiveValues<Integer> input = new ReactiveValues<>();
3     Integer r = 5;
4     ReactiveValues<Integer> output = new ReactiveValues<>();
5     output.subscribe();
6 }

```

```

1 (for ary(arity))
2   n1 ← Flt16toget(
3     testinput("n1", n1, "N1"),
4     testoutput("n1", n1)
5   )

```

**isolate(expr)**  
Runs a code block.  
Returns a **non-reactive**  
copy of the results.

```
function (input, output){
  output <-
    renderText({
      toString(input[i])
    })
}

shinyApp(ui, server)
```

```
var c = Reactive({
  transform: (v) => {
    (c.set, c.get) = (v, v);
  }
});

server.on('request', async() => {
  const {request, headers} =
    await c.request({url: 'http://localhost:3000/'});
});
```

[illegible]

```
library(sass)

ui <- function(
  textInput("a", "a"),
  textOutput("b"))

render*() functions
(see front page)
```

```

server <-
  function(input, output){
    on.TCPb <-
      renderText({
        input$b
      })
  }
shinyApp(ui, server)

```

```

14br ary(shiny)
us <- fluidPage(
  titlePanel("us"),
  actionButton("go", "Go")
)

```

**observeEvent(event,**  
**handlerExpr, env = \$R::**  
**environment(), hand arg an**  
**handler, closed, abn,**  
**and, and, critically dom**

```
server <
function (input, output) {
  observeEvent(input$a, {
    print(input$a)
  })
}

shinyApp(ui, server)
```

```

library(shiny)

# Create a reactive value
value <- reactive({
  actionButton("go", "Go")
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

```

```

library(shiny)

# Create a reactive value
value <- reactive({
  actionButton("go", "Go")
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

# Create a reactive expression with code in it
output <- renderText({
  value()
})

```

Use Shiny's functions to assemble this HTML with R.

[illegible][illegible][illegible]

```

    p(" ");
    b(" ");
    p("strong('bold')");
    p("se('cat','c')");
    p("code('code')");

```

|  |                            |
|--|----------------------------|
| $\frac{d(\ln \rho / \rho_0) / dt}{\ln \rho / \rho_0} = \frac{m}{\ln \rho / \rho_0} \left( \frac{d \ln \rho / \rho_0}{dt} \right)$ $\ln \rho / \rho_0 = \left( \frac{m}{\ln \rho / \rho_0} \right) \ln \rho / \rho_0$ | (for<br>$m =$<br>from item |
|--|----------------------------|

```
tags$head(tags$link(rel = "stylesheet",
type = "text/css", href = "file name"))
```

1. Place the file in the [www](#) subdirectory
2. Link to it with

2. Link to it with `img(src="<file name>")`

2. Link to it with `img(src="<file name>")`

Combine multiple elements into a "single element" that

```

    mtfPanel()
    confPanel()
    addPanel()
    delPanel()
}

```



Organize panels and elements into a layout with a layout function. Add elements as arguments of the layout function.

```

j1 <- 0, ifPage {
  fudgeBox(column1,width = 4,
    column1,width = 2, offset = 3),
  fudgeBox(column1,width = 12))

```

```

NewLayout: object 1,
#object: 2,
#object: 3

```

```
sideBarLayout()
ul < .fl {clearPage!
  sideBarLayout()
  sideBarParent(),
  mainParent()
}
```

```

spRLayout()
object
) | <- R.idPage()
spRLayout() # object 1,
# object 2

```

**verticalLayout()**

Layer tabPanels on top of each other, and navigate between them, with:

```

14 <!-- Page Navigation -->
15 <libPanel "no 1, contents"
16 <libPanel "no 2, contents"
17 <libPanel "no 3, contents"
18 </libPageNavigation>

```



# 14.4 R MARKDOWN cheat sheet [source: rstudio.com]

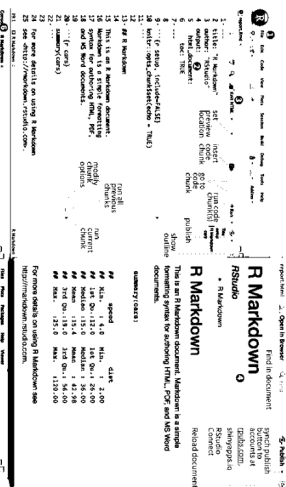


## R Markdown :: CHEAT SHEET

### What is R Markdown?

R Markdown is a format for creating documents that combine R code and prose. It is a way to create documents that are both reproducible and easy to read. The documents are created using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line. The documents can be created using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

### Workflow



1. Open a new R Markdown file in R Studio.
2. Write R code in the R code chunk.
3. Write prose in the prose chunk.
4. Render the document.
5. Export the document to PDF.

### Render

Render the document. The document is rendered into HTML or PDF format. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

### GLOBAL OPTIONS

GLOBAL OPTIONS. The document is rendered using the global options. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

GLOBAL OPTIONS. The document is rendered using the global options. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

GLOBAL OPTIONS. The document is rendered using the global options. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

Rstudio® is a trademark of RStudio, Inc. © 2015 RStudio, Inc. All rights reserved. RStudio is a trademark of RStudio, Inc. © 2015 RStudio, Inc. All rights reserved.

### Structure



### Parameters

Parameters. The document is rendered using the parameters. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

### Interactive Documents

Interactive Documents. The document is rendered using the interactive documents. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.

Interactive Documents. The document is rendered using the interactive documents. The document can be rendered using R Studio or the command line. The document can be rendered using R Markdown files, which are saved as .Rmd files. The documents are then rendered into HTML or PDF format. The documents can be created using R Studio or the command line.



